

MP4_P1

November 20, 2020

1 Generative Adversarial Networks

For this part of the assignment you implement two different types of generative adversarial networks. We will train the networks on the Celeb A dataset which is a large set of celebrity face images.

```
[ ]: import torch
from torch.utils.data import DataLoader
from torchvision import transforms
from torchvision.datasets import ImageFolder

import numpy as np

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

%load_ext autoreload
%autoreload 2

[ ]: from gan.train import train
[ ]: device = torch.device("cuda:0" if torch.cuda.is_available() else "gpu")
```

2 GAN loss functions

In this assignment you will implement two different types of GAN cost functions. You will first implement the loss from the [original GAN paper](#). You will also implement the loss from [LS-GAN](#).

```
[ ]: from gan.losses import discriminator_loss, generator_loss
```

2.0.1 Least Squares GAN loss

TODO: Implement the `ls_discriminator_loss` and `ls_generator_loss` functions in `gan/losses.py`.

We'll now look at [Least Squares GAN](#), a newer, more stable alternative to the original GAN loss function. For this part, all we have to do is change the loss function and retrain the model. We'll

implement equation (9) in the paper, with the generator loss:

$$\ell_G = \frac{1}{2} \mathbb{E}_{z \sim p(z)} [(D(G(z)) - 1)^2]$$

and the discriminator loss:

$$\ell_D = \frac{1}{2} \mathbb{E}_{x \sim p_{\text{data}}} [(D(x) - 1)^2] + \frac{1}{2} \mathbb{E}_{z \sim p(z)} [(D(G(z)))^2]$$

HINTS: Instead of computing the expectation, we will be averaging over elements of the mini-batch, so make sure to combine the loss by averaging instead of summing. When plugging in for $D(x)$ and $D(G(z))$ use the direct output from the discriminator (`scores_real` and `scores_fake`).

```
[ ]: from gan.losses import ls_discriminator_loss, ls_generator_loss
```

3 GAN model architecture

TODO: Implement the Discriminator and Generator networks in `gan/models.py`.

We recommend the following architectures which are inspired by [DCGAN](#):

Discriminator:

- convolutional layer with `in_channels=3, out_channels=128, kernel=4, stride=2`
- convolutional layer with `in_channels=128, out_channels=256, kernel=4, stride=2`
- batch norm
- convolutional layer with `in_channels=256, out_channels=512, kernel=4, stride=2`
- batch norm
- convolutional layer with `in_channels=512, out_channels=1024, kernel=4, stride=2`
- batch norm
- convolutional layer with `in_channels=1024, out_channels=1, kernel=4, stride=1`

Use padding = 1 (not 0) for all the convolutional layers.

Instead of Relu we LeakyReLu throughout the discriminator (we use a negative slope value of 0.2). You can use simply use relu as well.

The output of your discriminator should be a single value score corresponding to each input sample. See `torch.nn.LeakyReLU`.

Generator:

Note: In the generator, you will need to use transposed convolution (sometimes known as fractionally-strided convolution or deconvolution). This function is implemented in pytorch as `torch.nn.ConvTranspose2d`.

- transpose convolution with `in_channels=NOISE_DIM, out_channels=1024, kernel=4, stride=1`
- batch norm
- transpose convolution with `in_channels=1024, out_channels=512, kernel=4, stride=2`
- batch norm
- transpose convolution with `in_channels=512, out_channels=256, kernel=4, stride=2`
- batch norm
- transpose convolution with `in_channels=256, out_channels=128, kernel=4, stride=2`
- batch norm

- transpose convolution with in_channels=128, out_channels=3, kernel=4, stride=2

The output of the final layer of the generator network should have a `tanh` nonlinearity to output values between -1 and 1. The output should be a 3x64x64 tensor for each sample (equal dimensions to the images from the dataset).

```
[ ]: from gan.models import Discriminator, Generator
```

4 Data loading: Celeb A Dataset

The CelebA images we provide have been filtered to obtain only images with clear faces and have been cropped and downsampled to 128x128 resolution.

```
[ ]: batch_size = 128
      scale_size = 64 # We resize the images to 64x64 for training

      celeba_root = 'celeba_data'

[ ]: celeba_train = ImageFolder(root=celeba_root, transform=transforms.Compose([
      transforms.Resize(scale_size),
      transforms.ToTensor(),
    ]))

      celeba_loader_train = DataLoader(celeba_train, batch_size=batch_size, drop_last=True)
```

4.0.1 Visualize dataset

```
[ ]: imgs = celeba_loader_train.__iter__().next()[0].numpy().squeeze()
      show_images(imgs, color=True)
```

5 Training

TODO: Fill in the training loop in `gan/train.py`.

```
[ ]: NOISE_DIM = 100
      NUM_EPOCHS = 50
      learning_rate = 0.0002
```

5.0.1 Train GAN

```
[ ]: D = Discriminator().to(device)
      G = Generator(noise_dim=NOISE_DIM).to(device)

[ ]: D_optimizer = torch.optim.Adam(D.parameters(), lr=learning_rate, betas = (0.5, 0.999))
      G_optimizer = torch.optim.Adam(G.parameters(), lr=learning_rate, betas = (0.5, 0.999))
```

```
[ ]: # original gan - spectral normalization
train(D, G, D_optimizer, G_optimizer, discriminator_loss,
      generator_loss, num_epochs=NUM_EPOCHS, show_every=150,
      train_loader=celeba_loader_train, device=device)
```

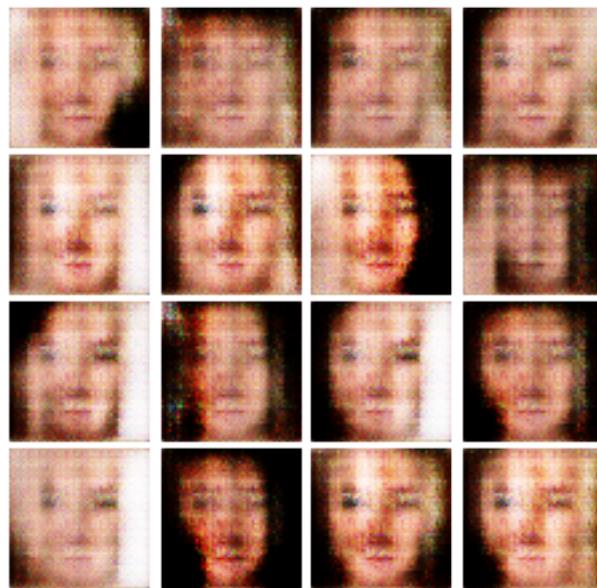
EPOCH: 1
Iter: 0, D: 1.464, G:5.203



Iter: 150, D: 0.4822, G:4.225



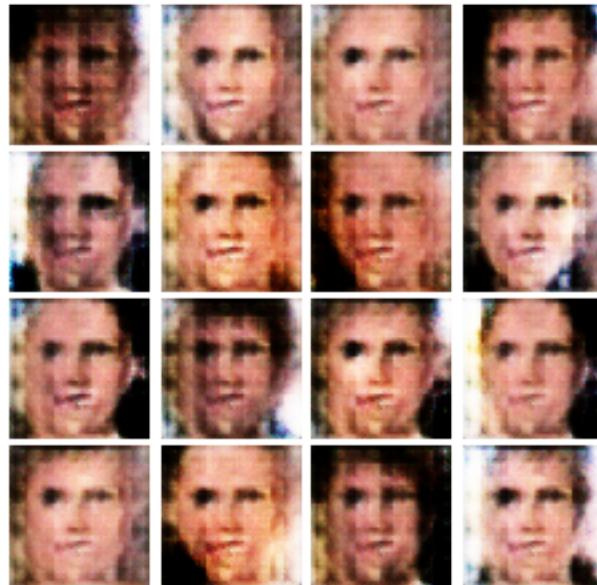
Iter: 300, D: 0.8303, G:5.985



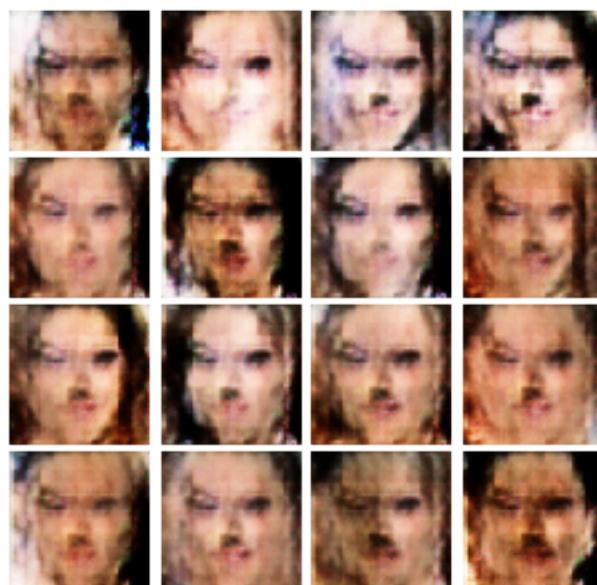
Iter: 450, D: 0.6145, G:2.894



Iter: 600, D: 1.239, G:2.7



Iter: 750, D: 0.6017, G:3.364



Iter: 900, D: 0.7745, G:3.329



EPOCH: 2

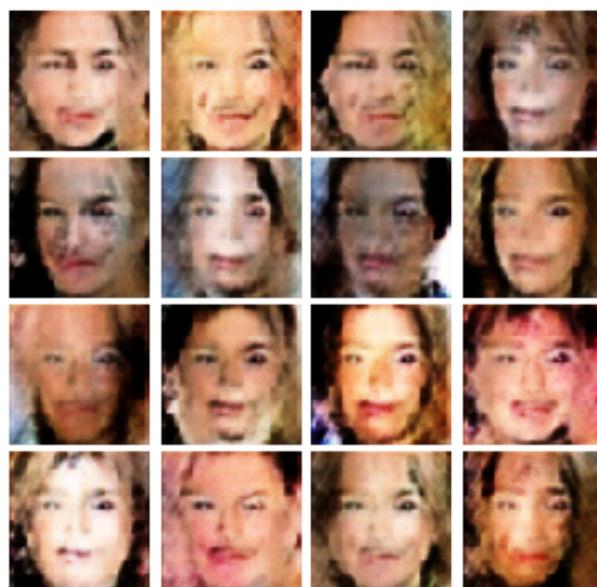
Iter: 1050, D: 0.787, G:3.696



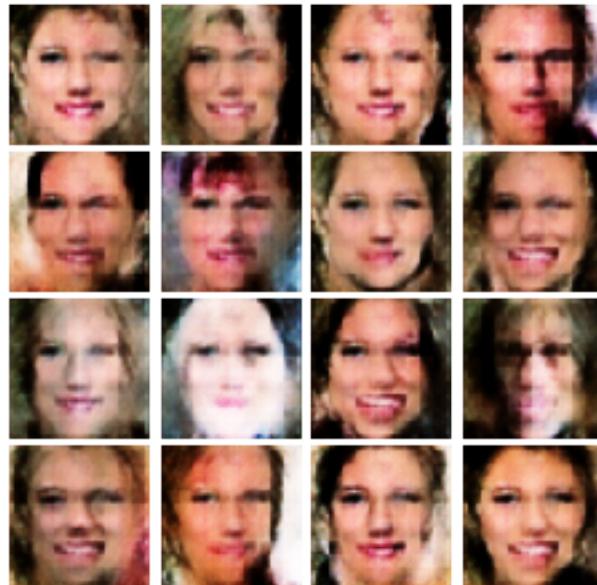
Iter: 1200, D: 0.6609, G:4.173



Iter: 1350, D: 0.4499, G:2.334



Iter: 1500, D: 1.141, G:1.866



Iter: 1650, D: 0.8206, G:4.874



Iter: 1800, D: 0.4146, G:3.103

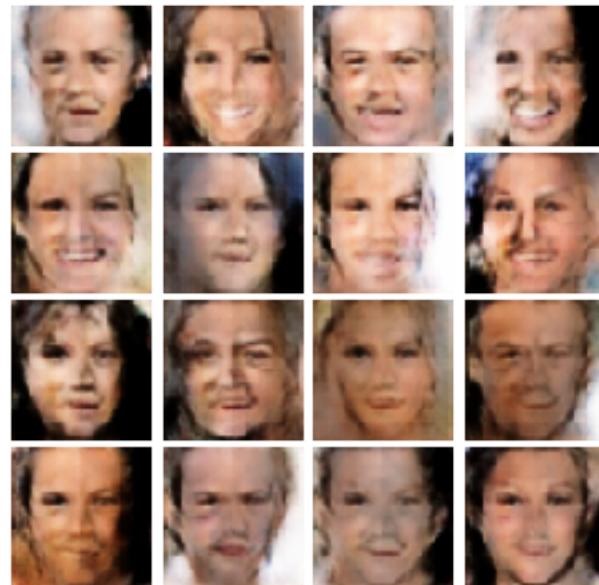


Iter: 1950, D: 0.2798, G:4.045

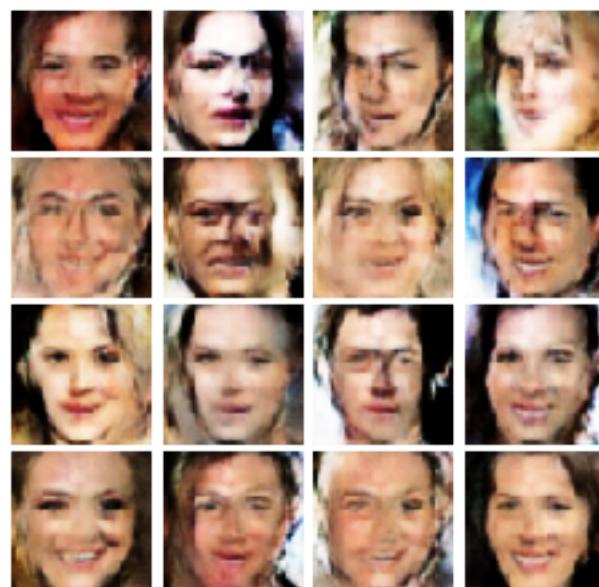


EPOCH: 3

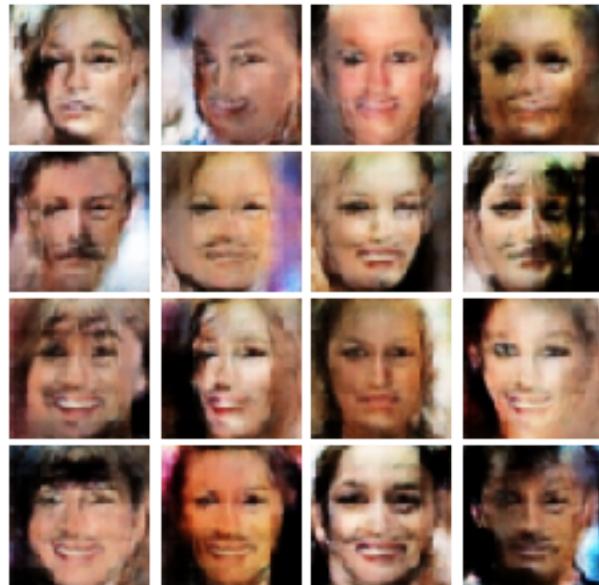
Iter: 2100, D: 0.4125, G:2.982



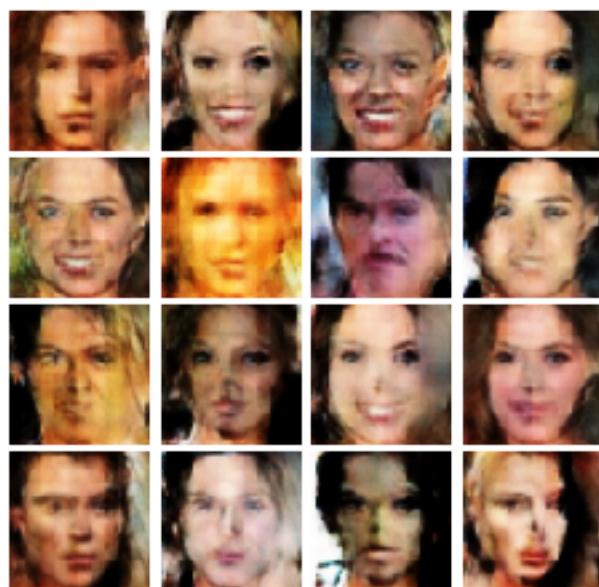
Iter: 2250, D: 0.4839, G:2.603



Iter: 2400, D: 0.4345, G:2.079



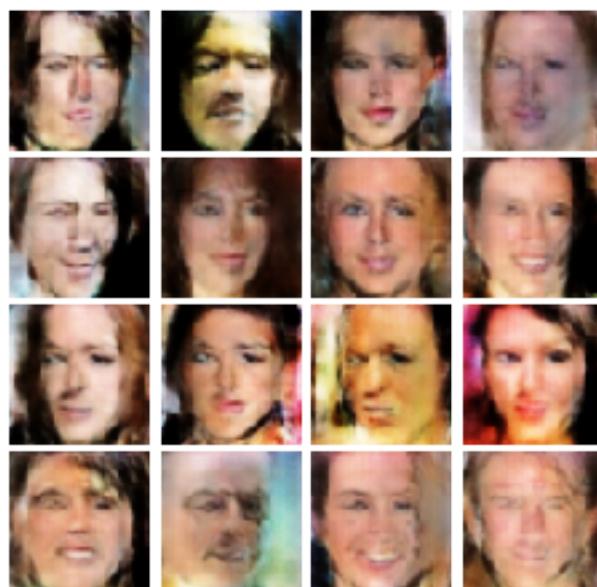
Iter: 2550, D: 0.2387, G:4.919



Iter: 2700, D: 0.2685, G:2.578

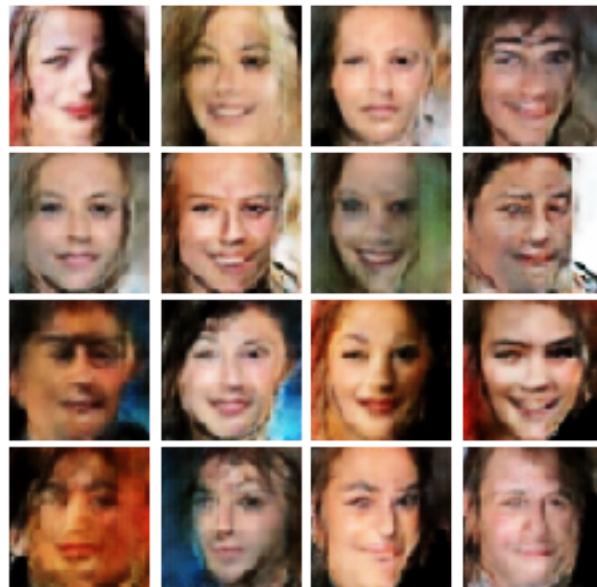


Iter: 2850, D: 0.473, G:2.788

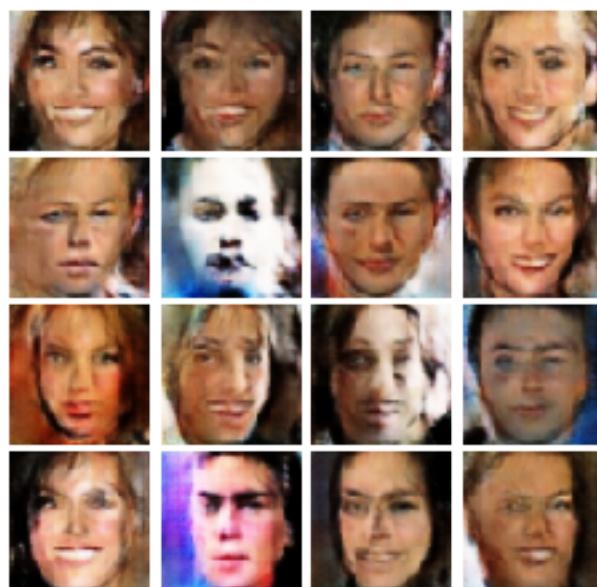


EPOCH: 4

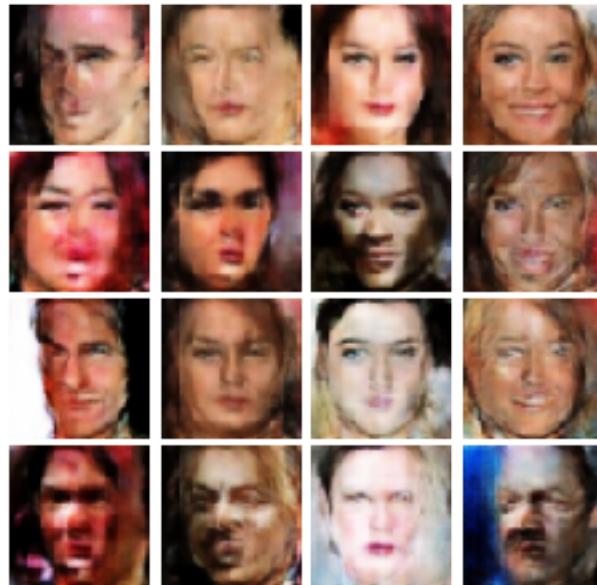
Iter: 3000, D: 0.6404, G:4.263



Iter: 3150, D: 0.1776, G:3.901



Iter: 3300, D: 0.5424, G:8.048



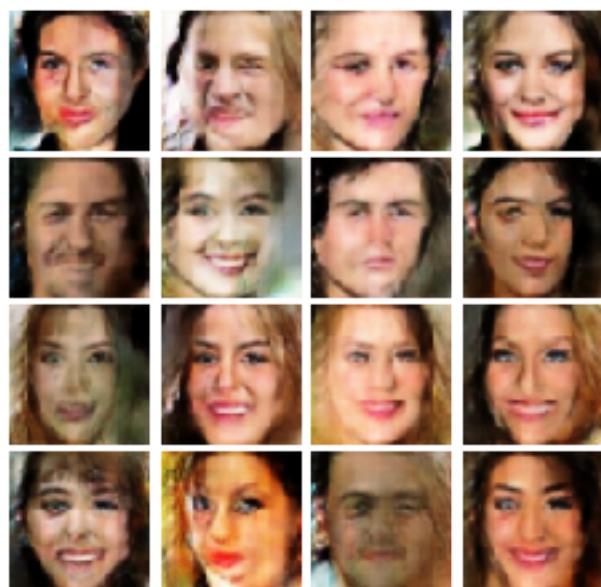
Iter: 3450, D: 1.133, G:2.46



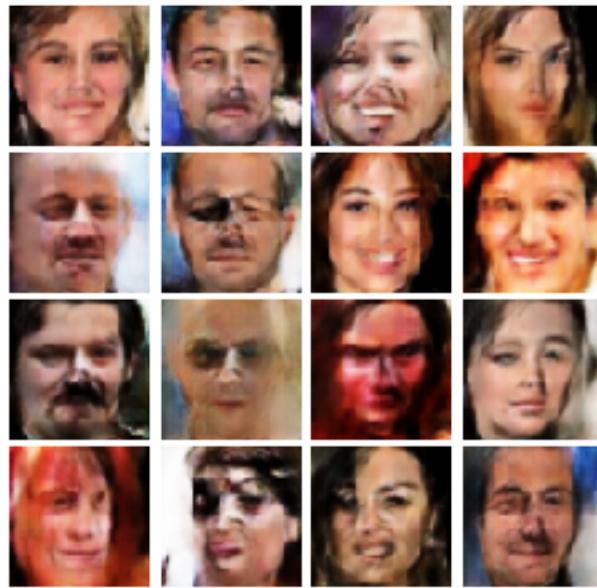
Iter: 3600, D: 1.615, G:15.39



Iter: 3750, D: 0.6859, G:1.75

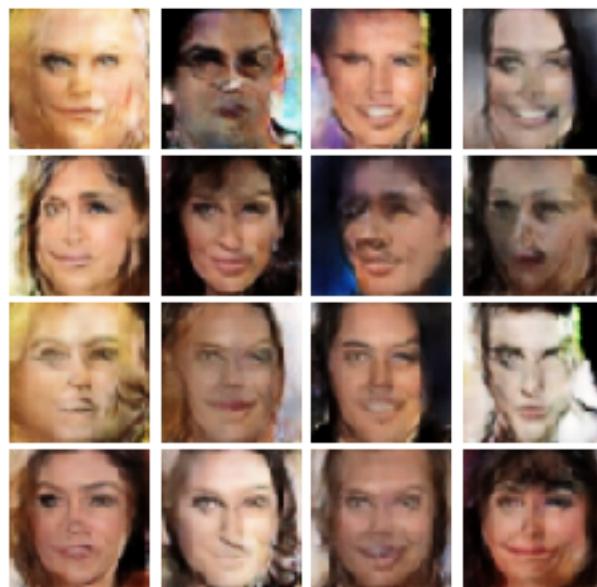


Iter: 3900, D: 0.5491, G:3.929



EPOCH: 5

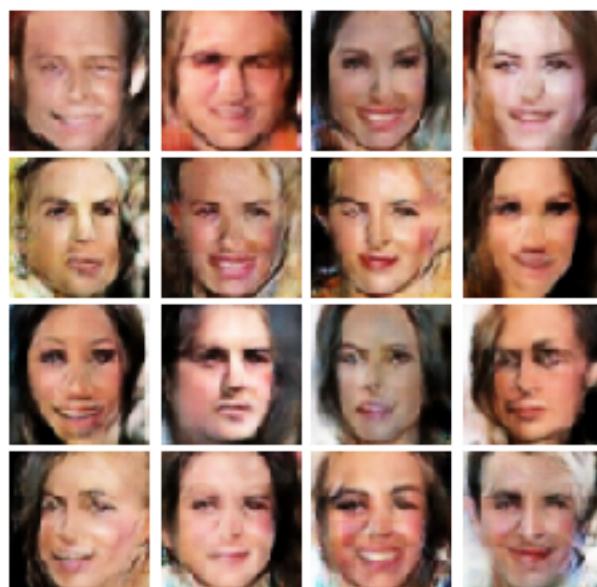
Iter: 4050, D: 1.137, G:2.713



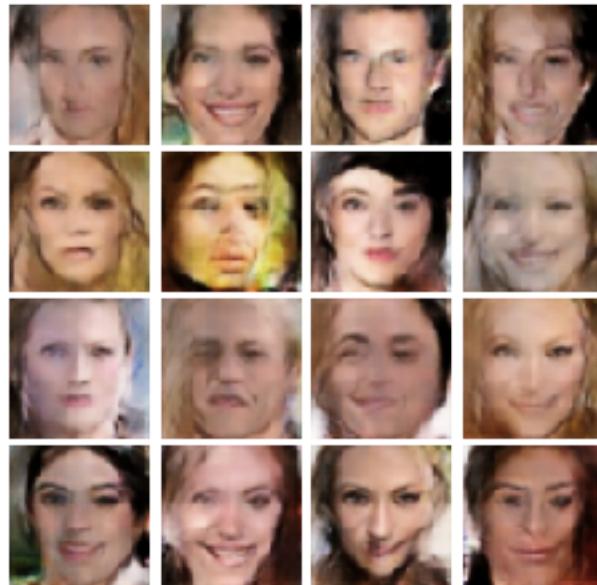
Iter: 4200, D: 0.1574, G:4.54



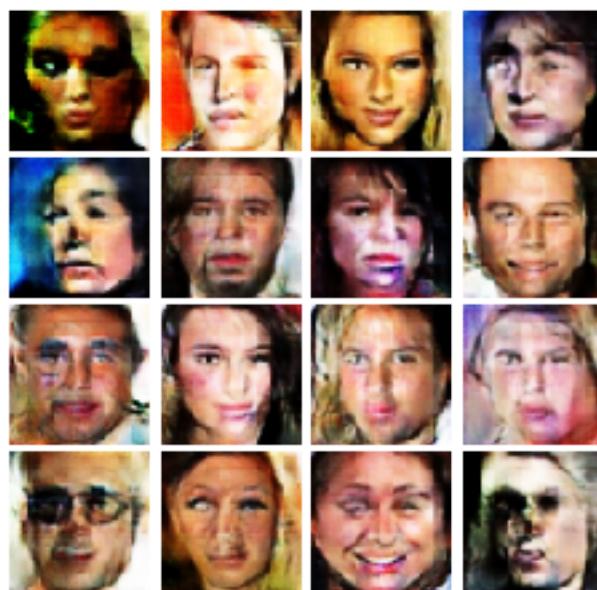
Iter: 4350, D: 0.7175, G:2.076



Iter: 4500, D: 0.2321, G:2.25



Iter: 4650, D: 0.1315, G:4.514

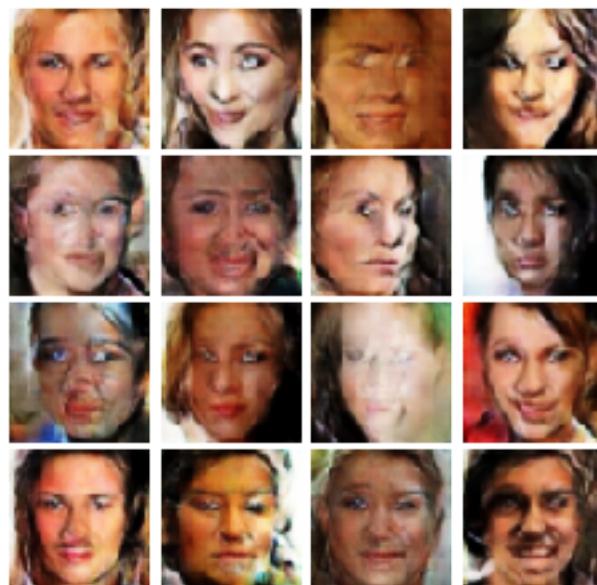


Iter: 4800, D: 0.2015, G:4.097

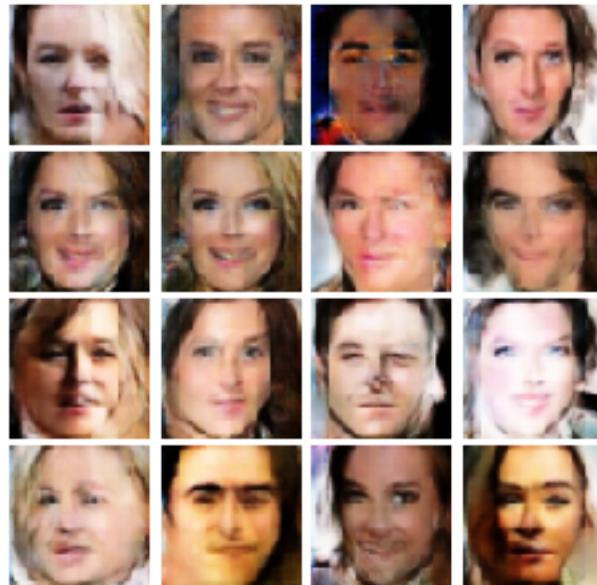


EPOCH: 6

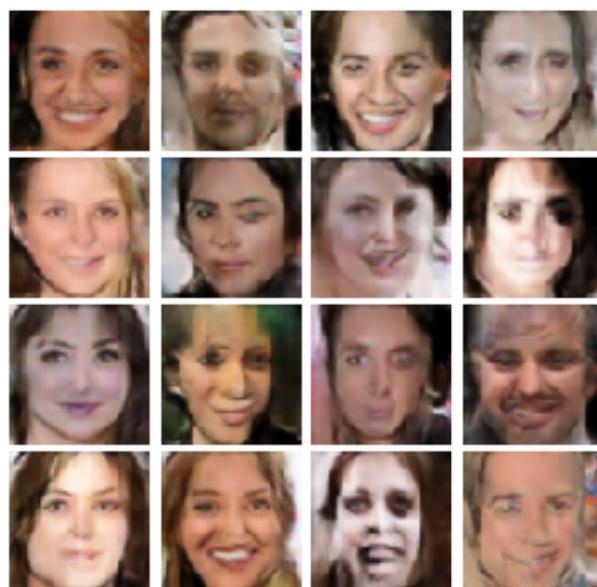
Iter: 4950, D: 0.5905, G:8.918



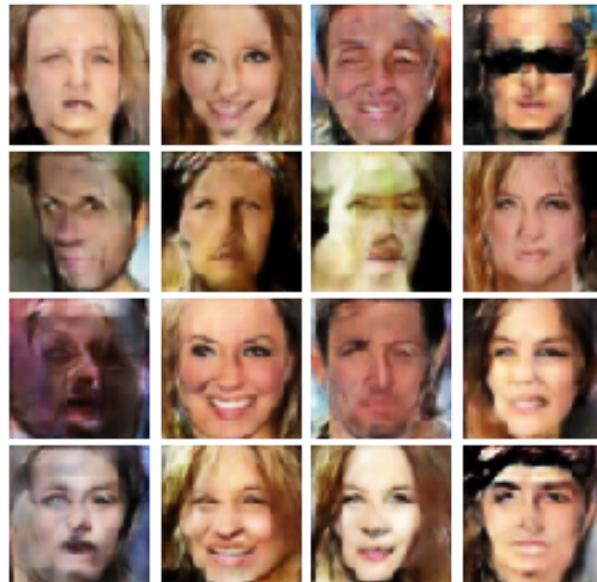
Iter: 5100, D: 1.165, G:7.783



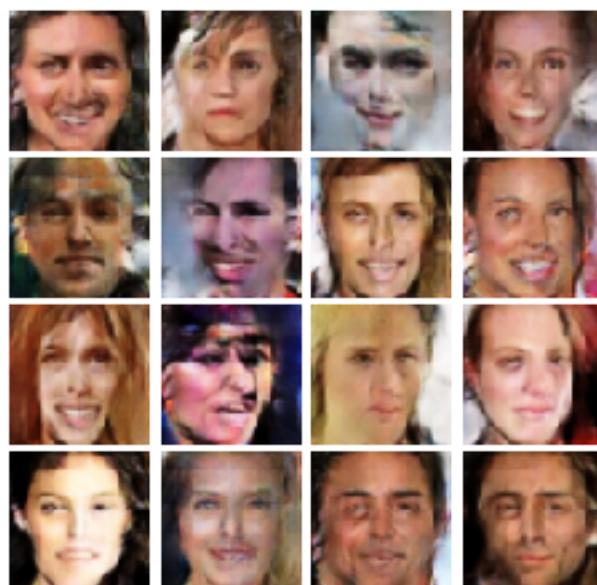
Iter: 5250, D: 2.397, G:11.24



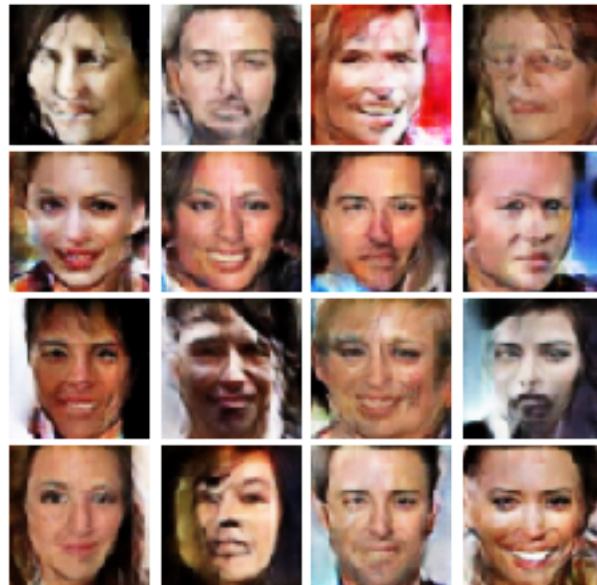
Iter: 5400, D: 0.2227, G:3.263



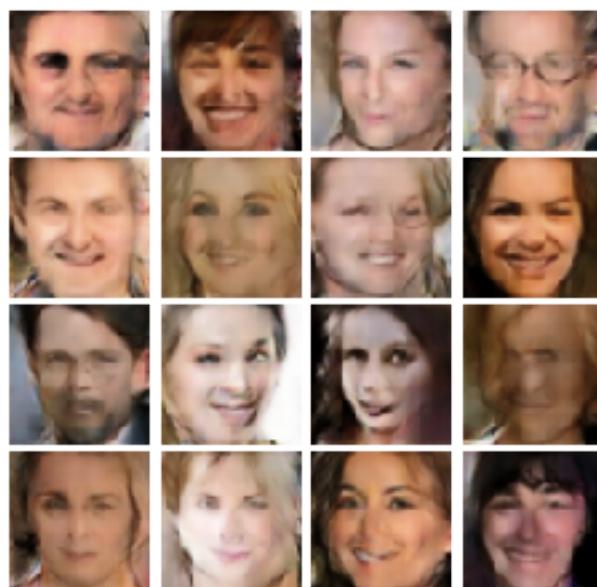
Iter: 5550, D: 0.224, G:3.16



Iter: 5700, D: 0.3269, G:3.169

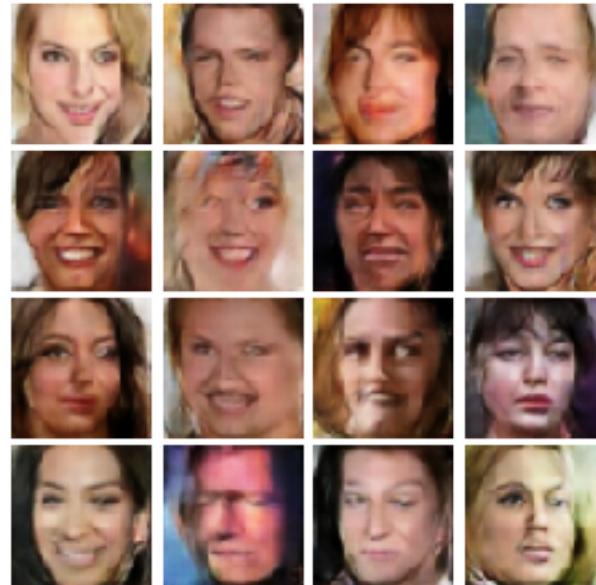


Iter: 5850, D: 0.3931, G:3.92

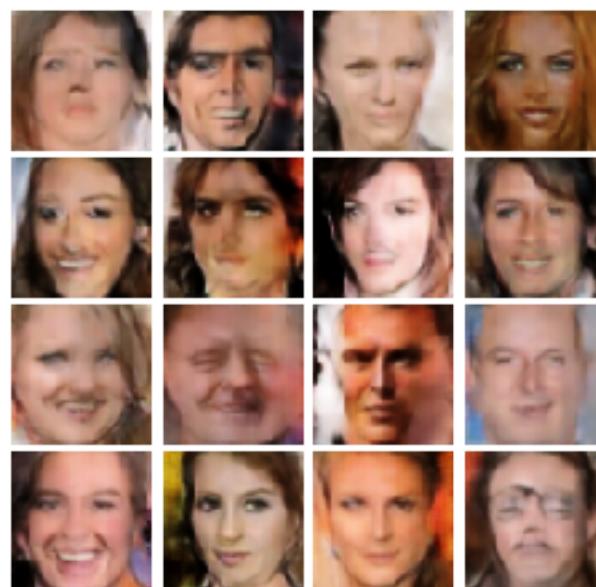


EPOCH: 7

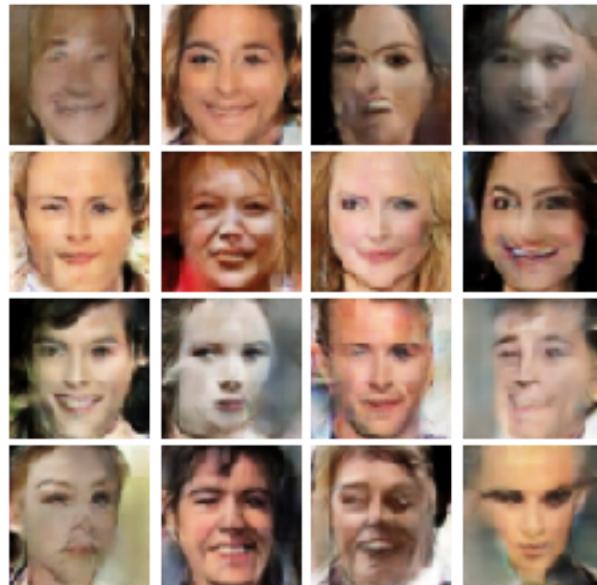
Iter: 6000, D: 0.1538, G:3.64



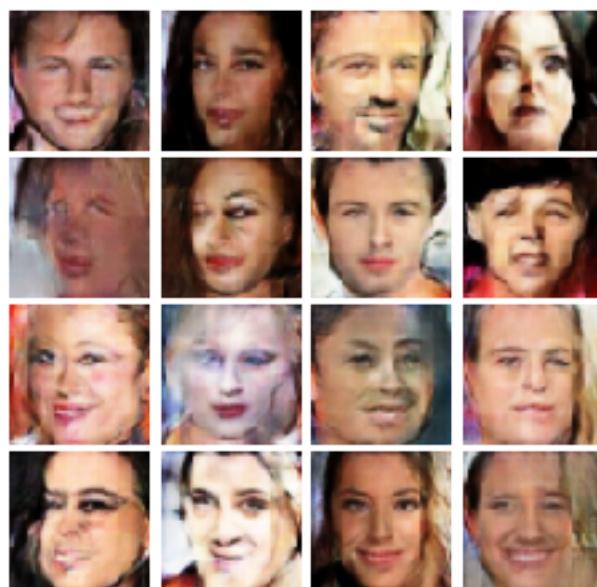
Iter: 6150, D: 0.7282, G:6.024



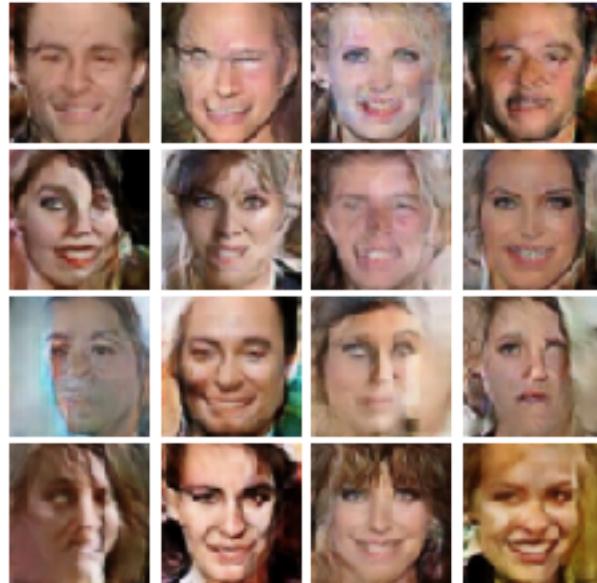
Iter: 6300, D: 0.1785, G:3.802



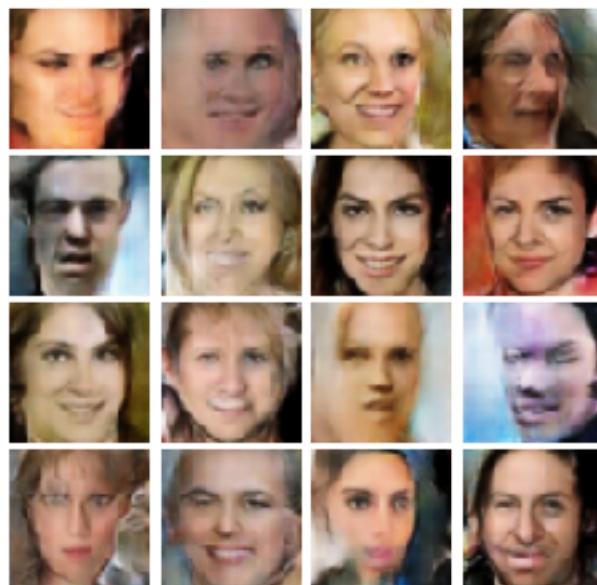
Iter: 6450, D: 0.124, G:2.671



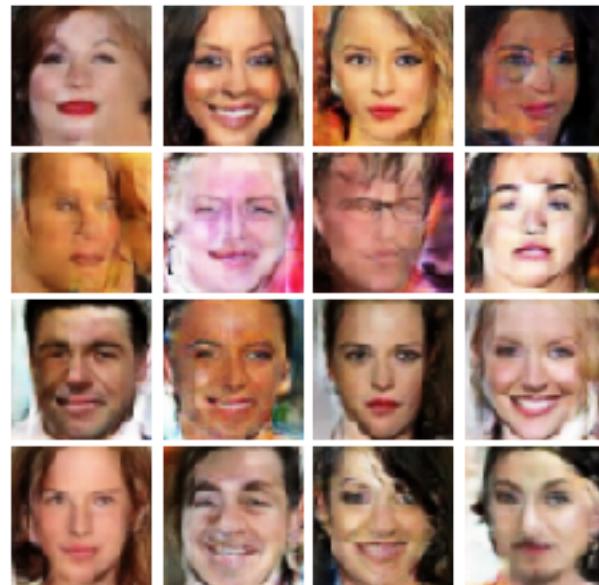
Iter: 6600, D: 0.2779, G:3.952



Iter: 6750, D: 0.3234, G:5.367

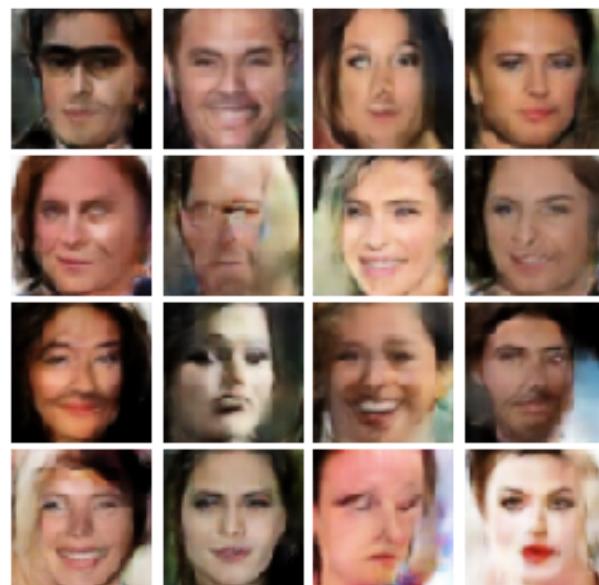


Iter: 6900, D: 0.1961, G:3.06

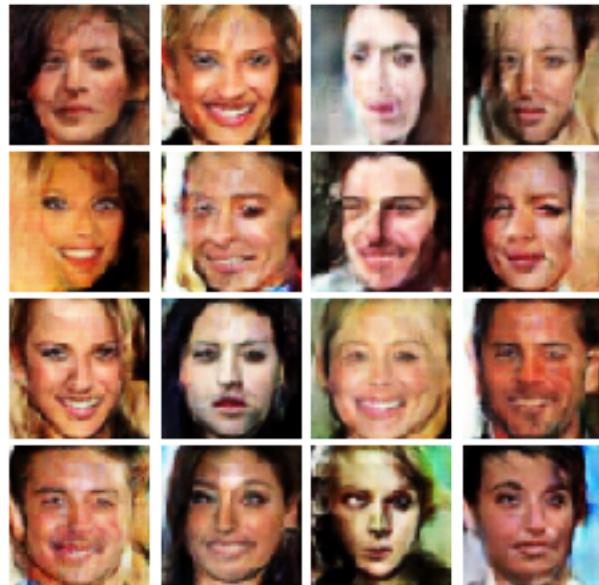


EPOCH: 8

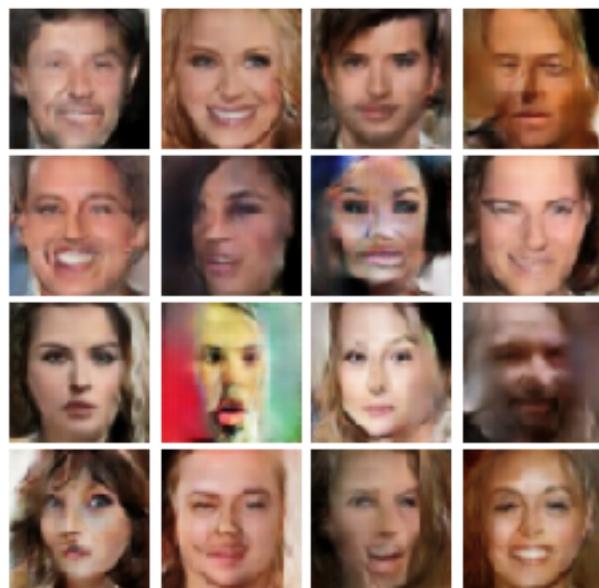
Iter: 7050, D: 0.08181, G:7.152



Iter: 7200, D: 0.3232, G:4.937



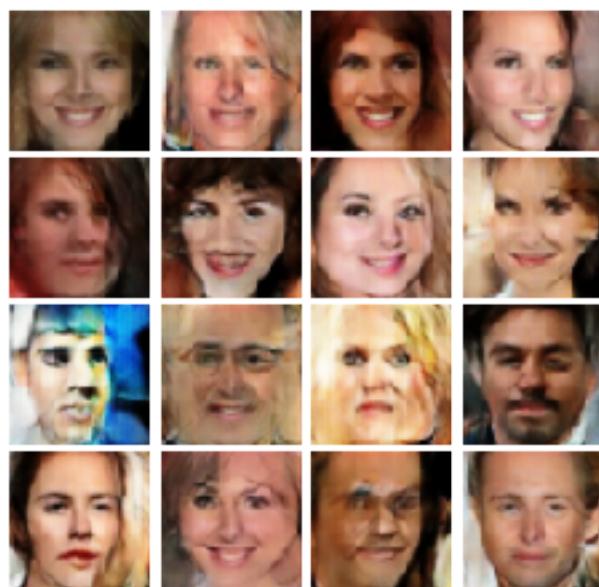
Iter: 7350, D: 0.3526, G:2.444



Iter: 7500, D: 0.05562, G:4.744



Iter: 7650, D: 0.1504, G:1.497

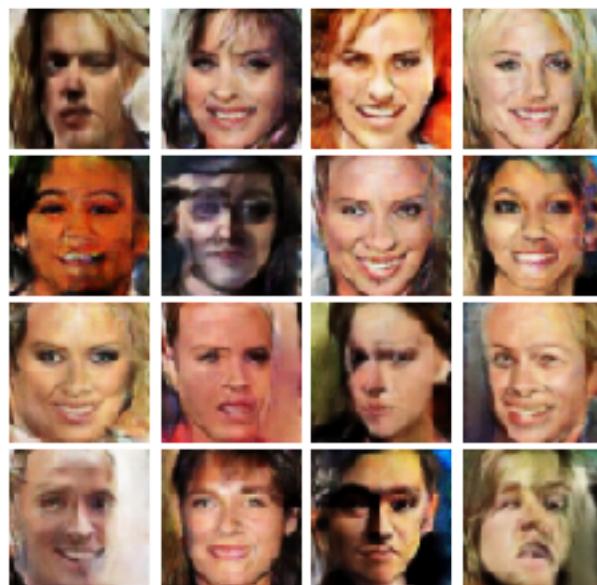


Iter: 7800, D: 0.04531, G:5.386

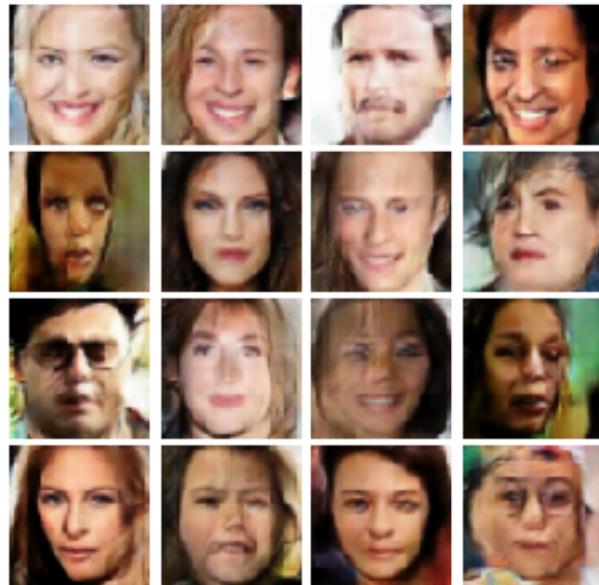


EPOCH: 9

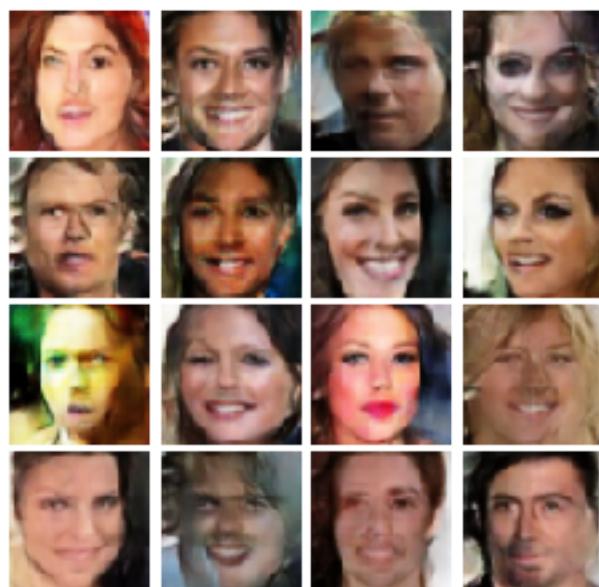
Iter: 7950, D: 0.8338, G:13.14



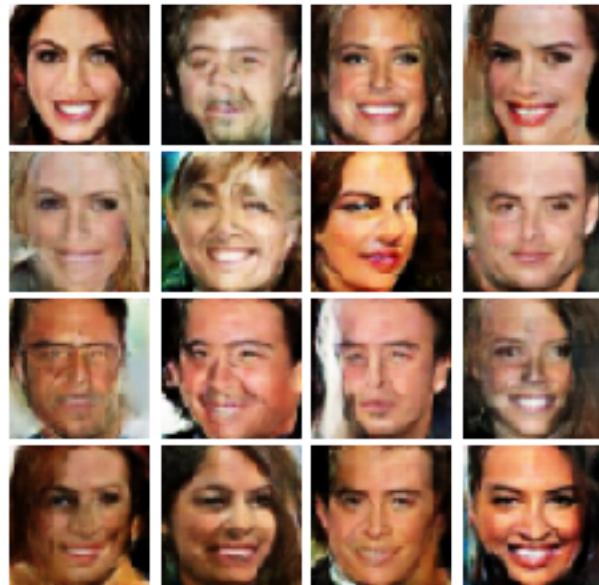
Iter: 8100, D: 0.1056, G:4.341



Iter: 8250, D: 0.08754, G:3.046



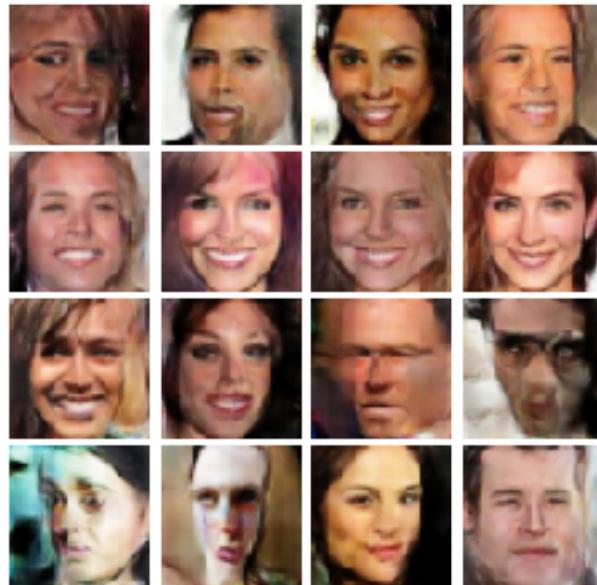
Iter: 8400, D: 0.1849, G:3.37



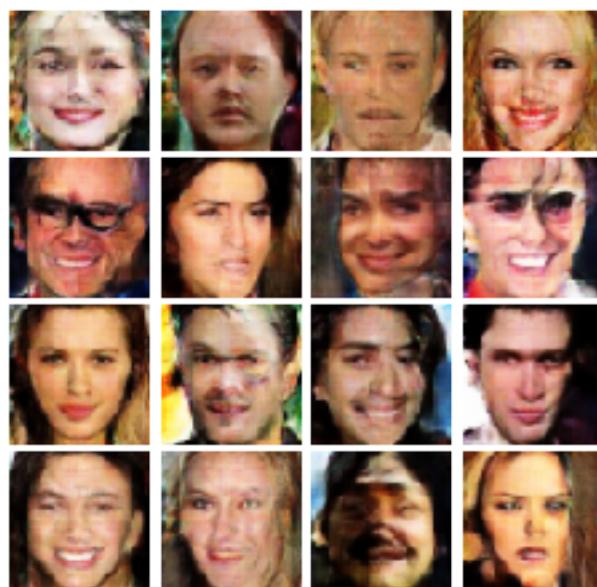
Iter: 8550, D: 0.2251, G:4.905



Iter: 8700, D: 0.7274, G:7.648

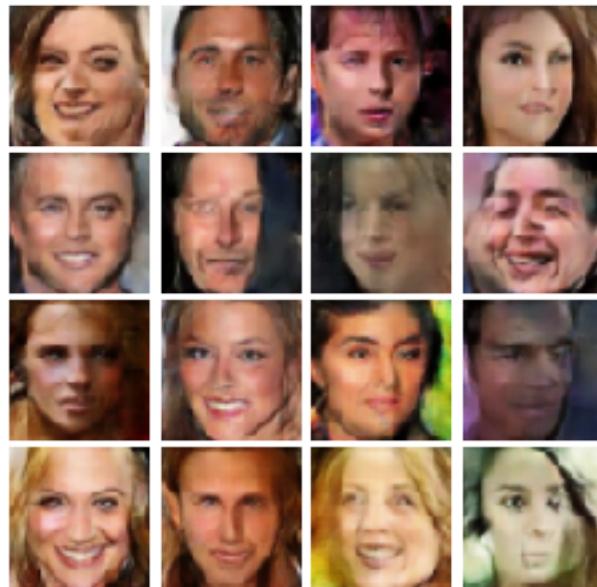


Iter: 8850, D: 0.106, G:4.847

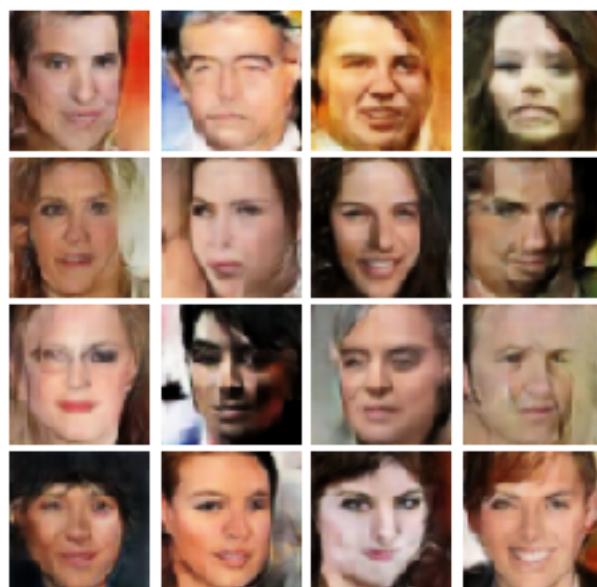


EPOCH: 10

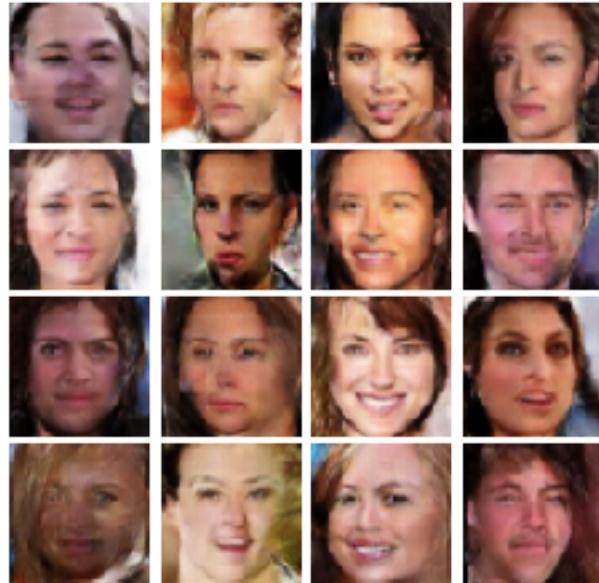
Iter: 9000, D: 0.04889, G:2.622



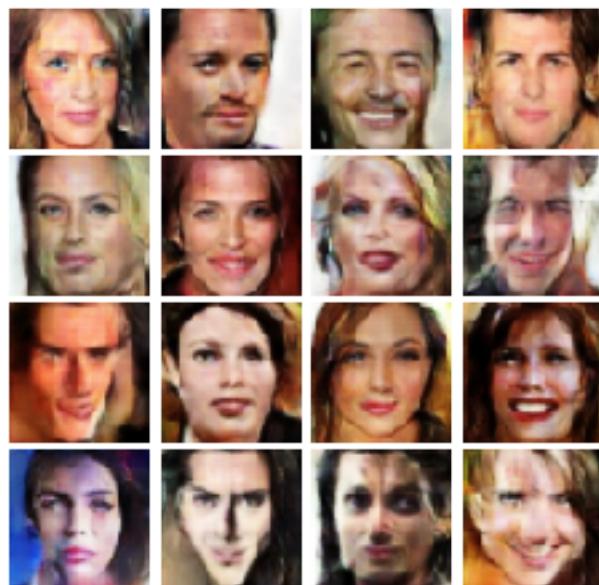
Iter: 9150, D: 0.2359, G:0.4249



Iter: 9300, D: 0.5293, G:4.905



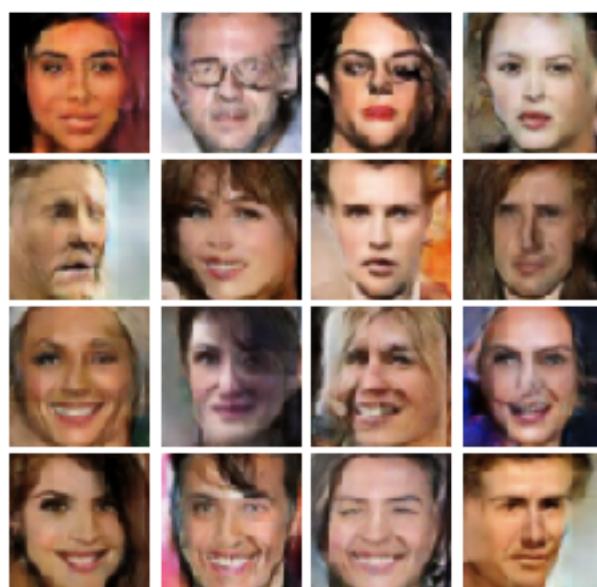
Iter: 9450, D: 0.1915, G:3.614



Iter: 9600, D: 0.3267, G:5.088

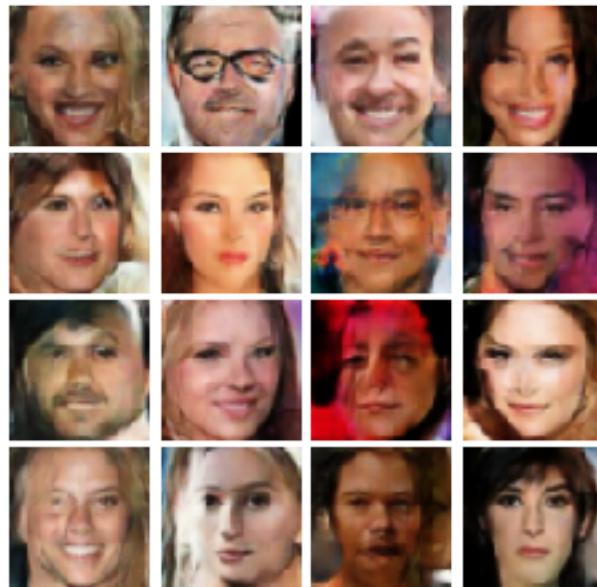


Iter: 9750, D: 0.09309, G:4.37

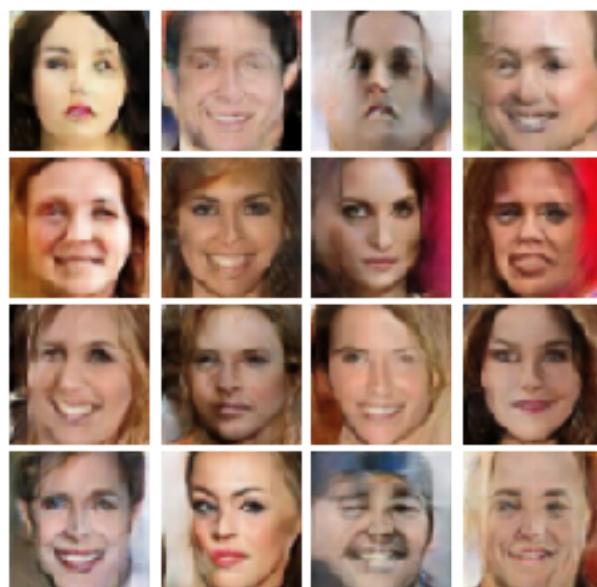


EPOCH: 11

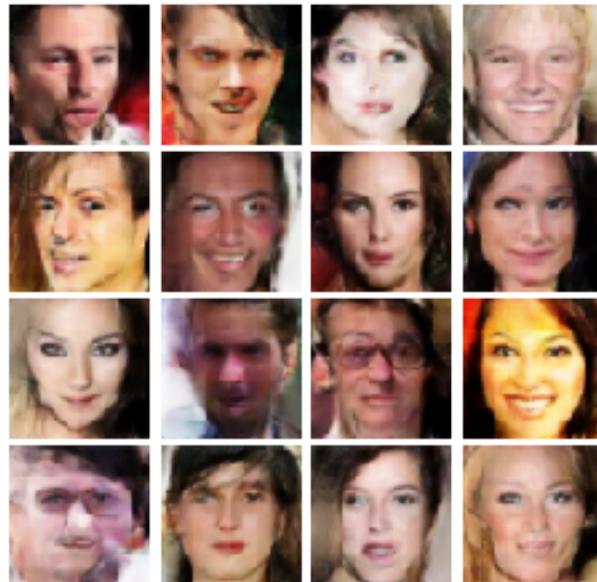
Iter: 9900, D: 0.07503, G:4.611



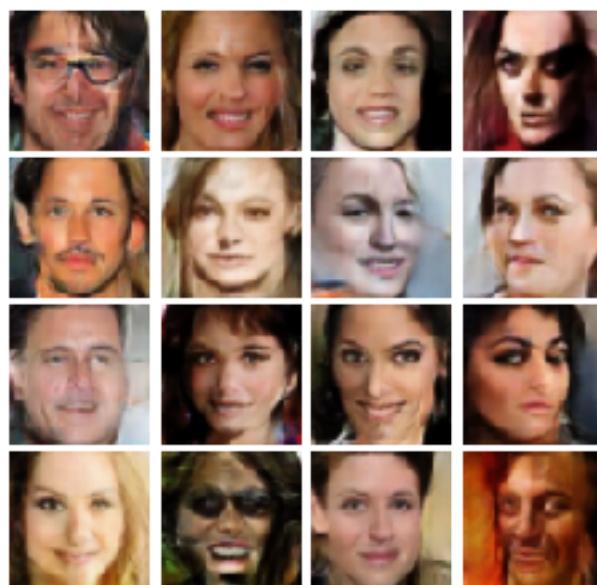
Iter: 10050, D: 0.1382, G:3.472



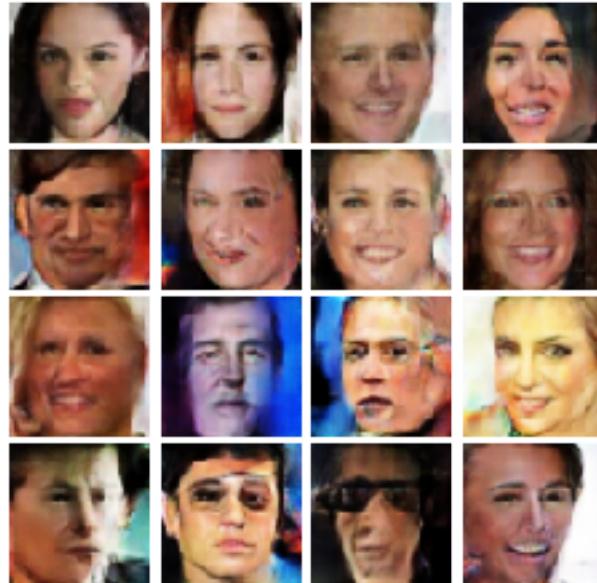
Iter: 10200, D: 0.8346, G:1.88



Iter: 10350, D: 0.2293, G:3.036



Iter: 10500, D: 0.3636, G:6.611



Iter: 10650, D: 0.2277, G:3.173

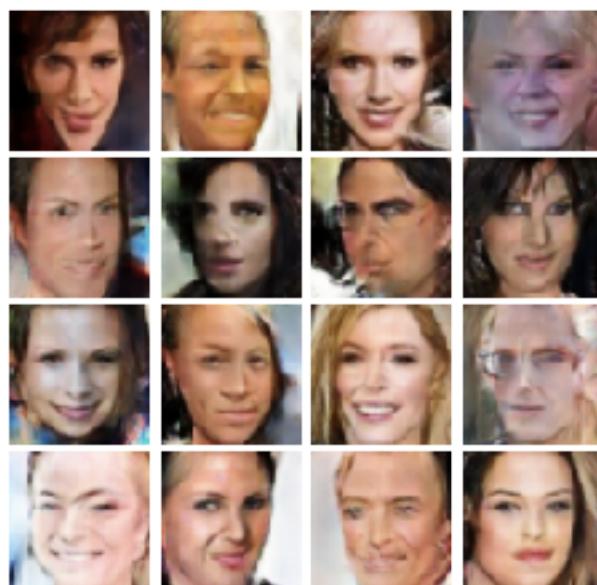


Iter: 10800, D: 0.2869, G: 5.153



EPOCH: 12

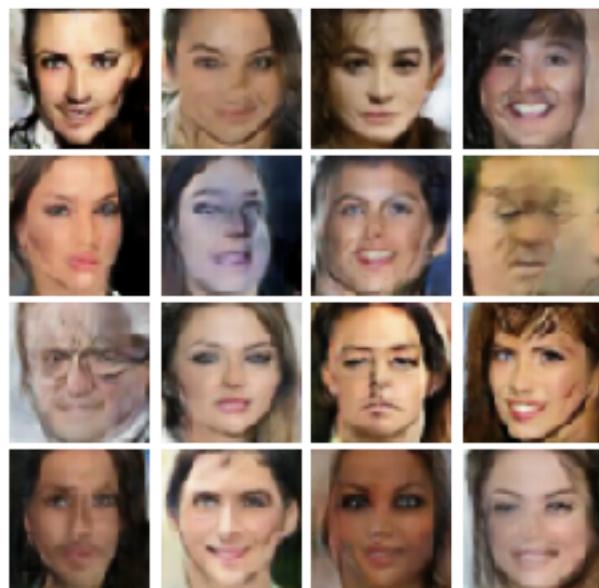
Iter: 10950, D: 0.08887, G: 2.443



Iter: 11100, D: 0.1534, G:5.224



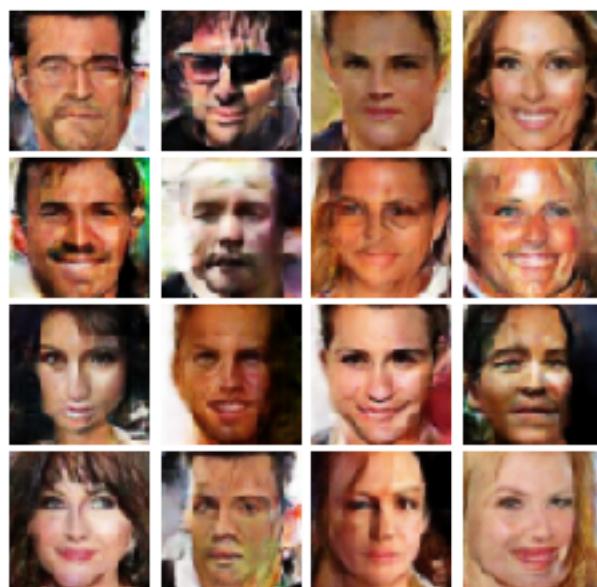
Iter: 11250, D: 0.2783, G:6.173



Iter: 11400, D: 0.1129, G:2.101



Iter: 11550, D: 0.07253, G:4.266

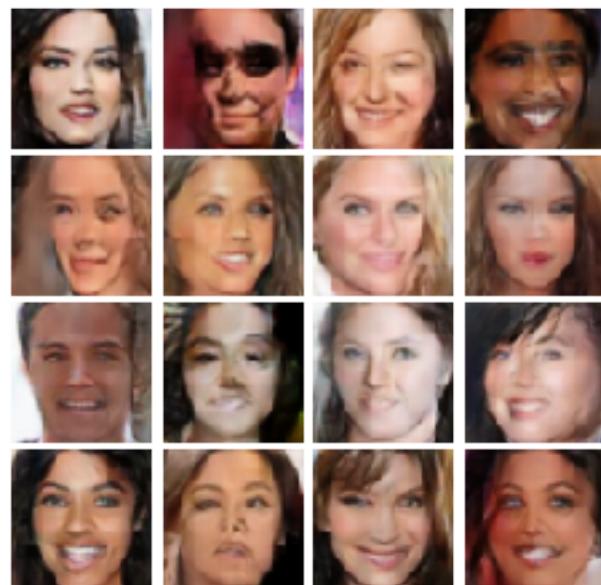


Iter: 11700, D: 0.03409, G:8.207

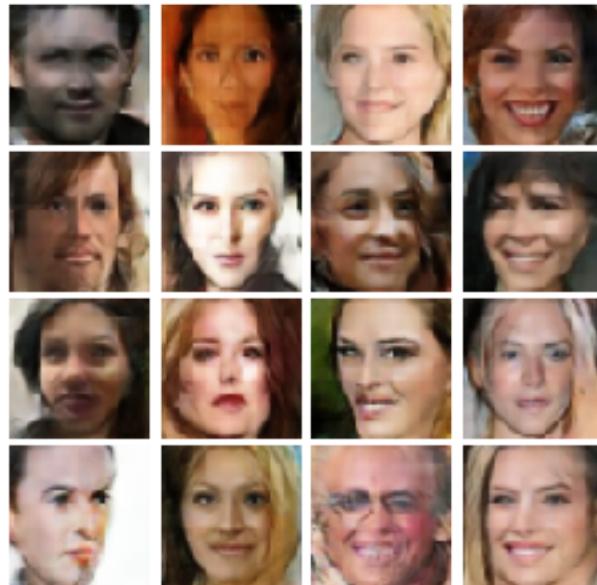


EPOCH: 13

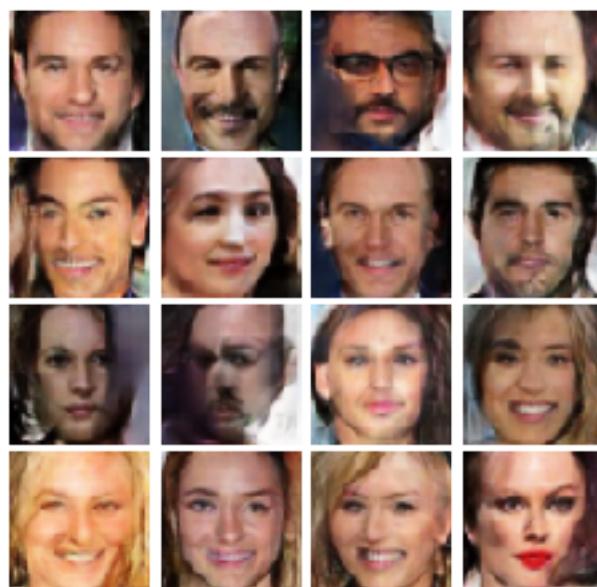
Iter: 11850, D: 0.1043, G:4.642



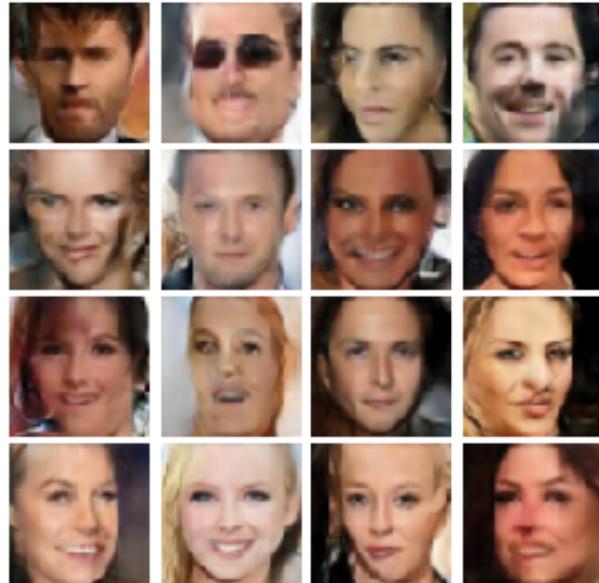
Iter: 12000, D: 0.05386, G:3.842



Iter: 12150, D: 0.213, G:5.146



Iter: 12300, D: 0.1097, G:5.455



Iter: 12450, D: 0.139, G:5.997



Iter: 12600, D: 0.2491, G:3.242

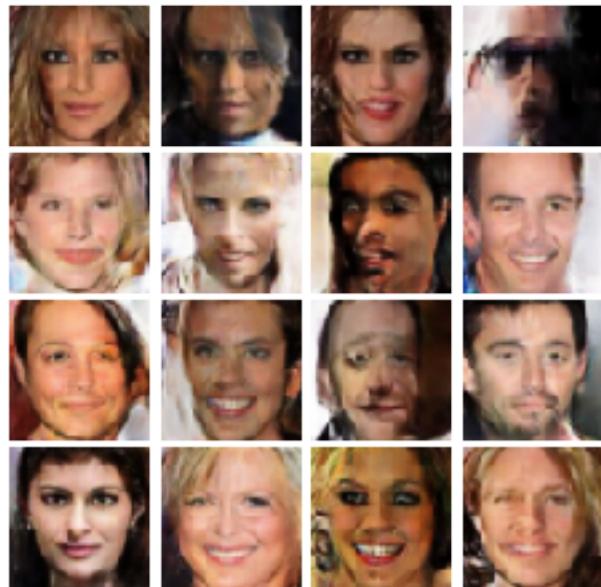


Iter: 12750, D: 0.06846, G:2.148

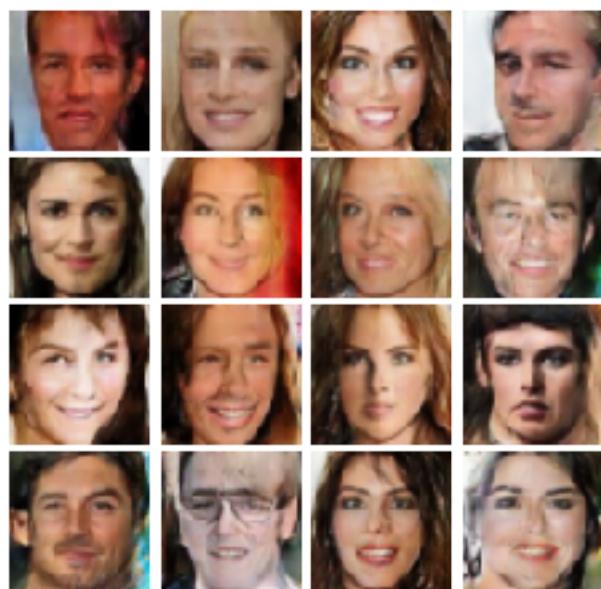


EPOCH: 14

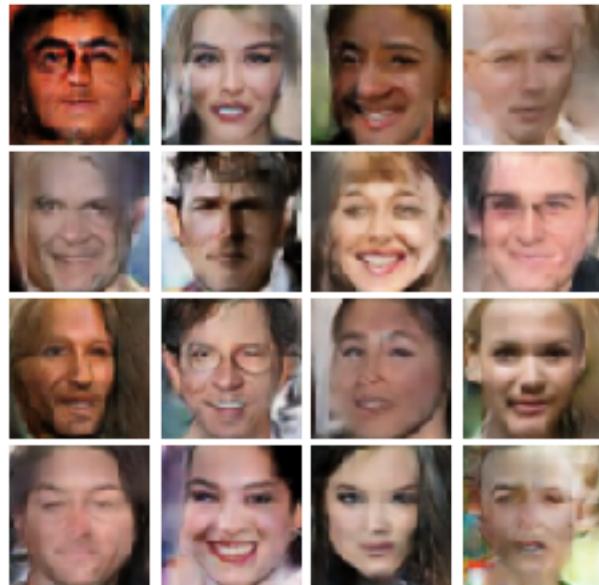
Iter: 12900, D: 1.691, G:1.045



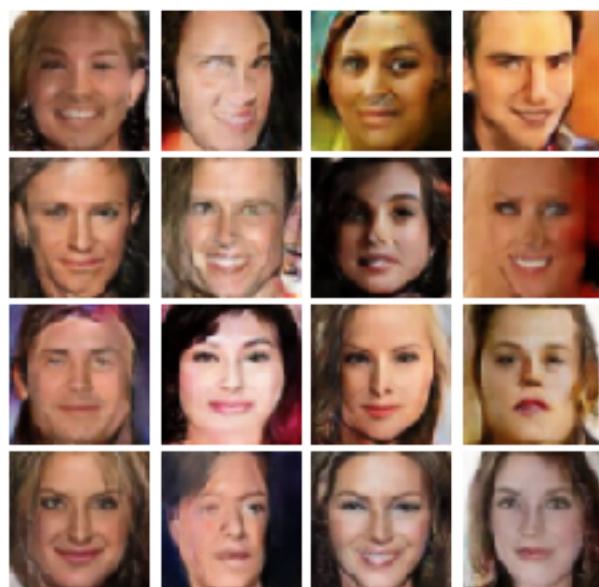
Iter: 13050, D: 0.1673, G:3.988



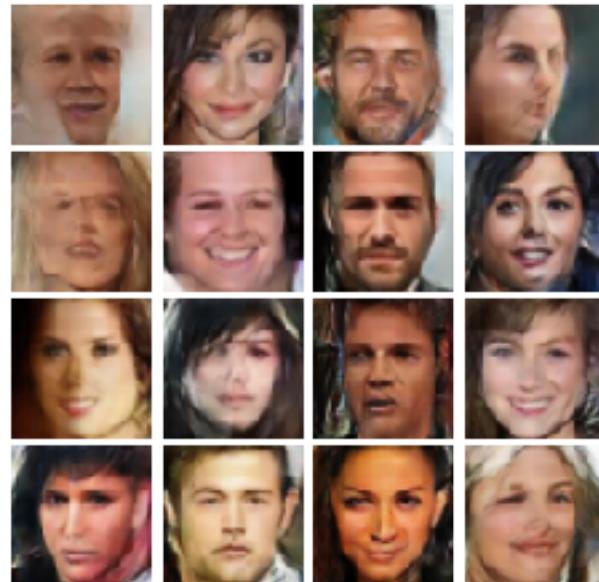
Iter: 13200, D: 0.7094, G:5.545



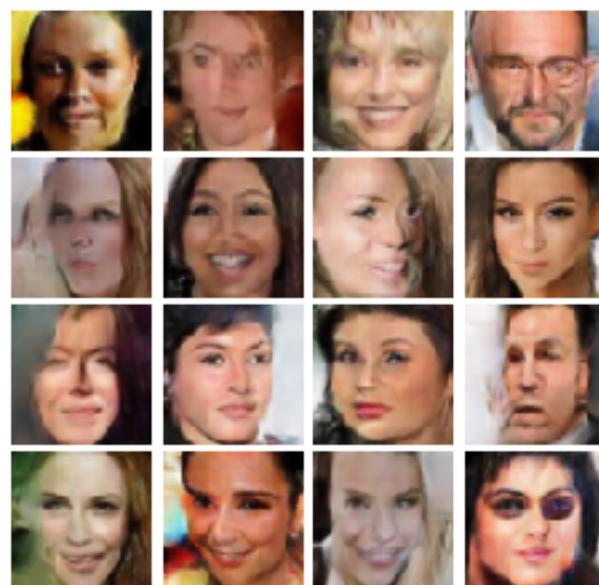
Iter: 13350, D: 0.103, G:4.158



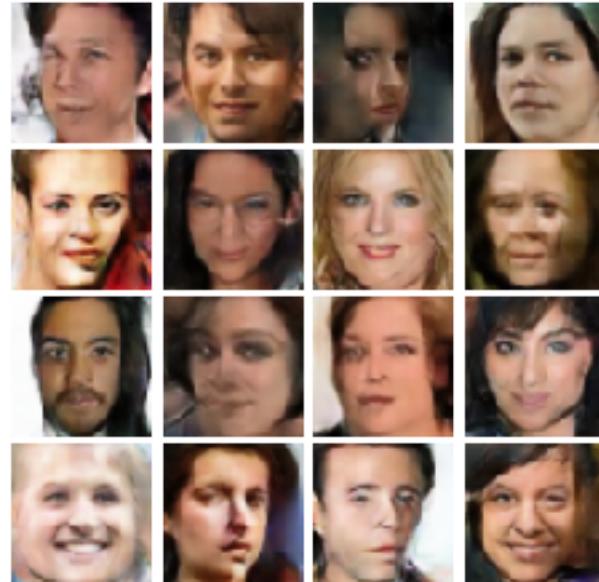
Iter: 13500, D: 0.08307, G:4.463



Iter: 13650, D: 0.05793, G:4.501

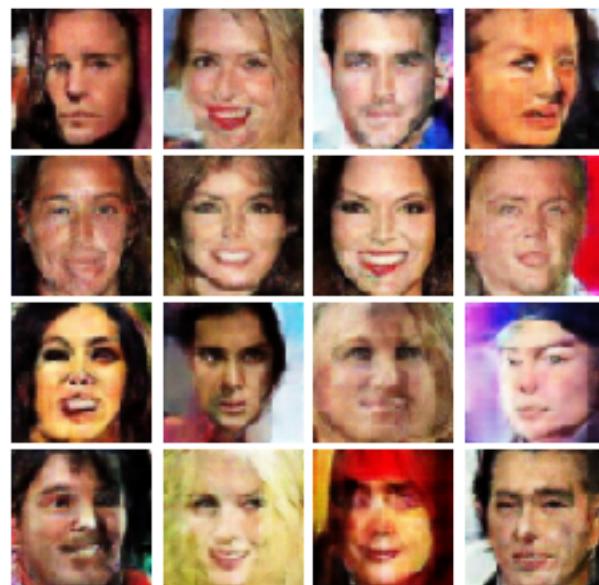


Iter: 13800, D: 0.02267, G:7.597

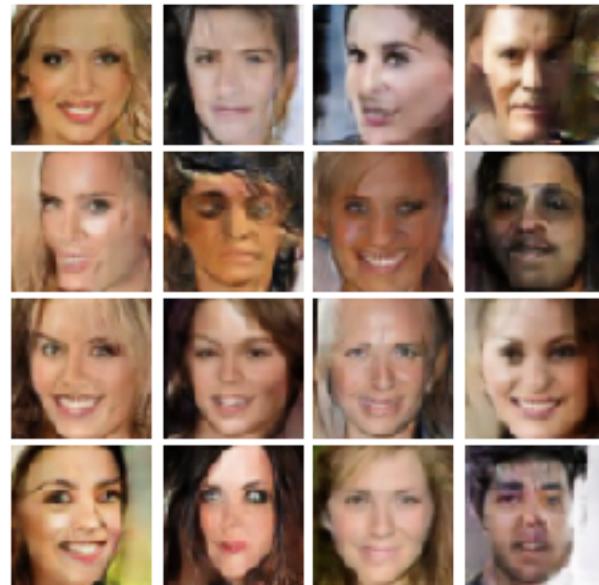


EPOCH: 15

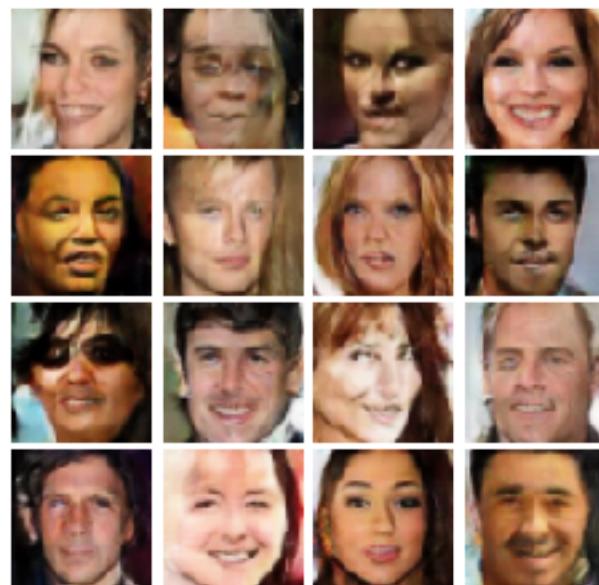
Iter: 13950, D: 0.1752, G:8.92



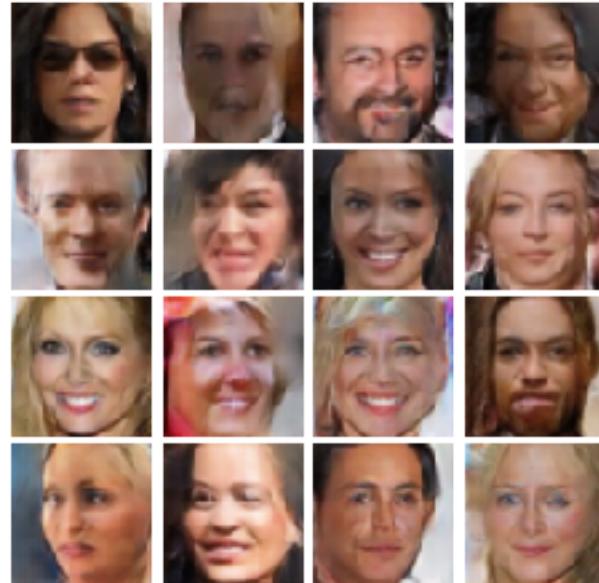
Iter: 14100, D: 0.1385, G:5.169



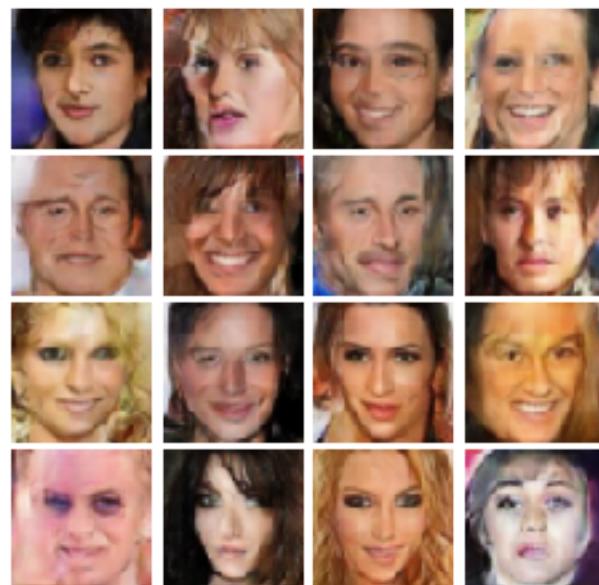
Iter: 14250, D: 0.3743, G:7.438



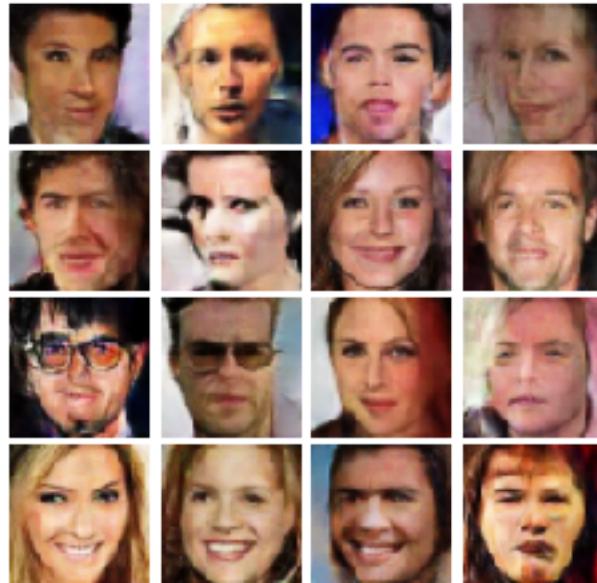
Iter: 14400, D: 0.4341, G:2.221



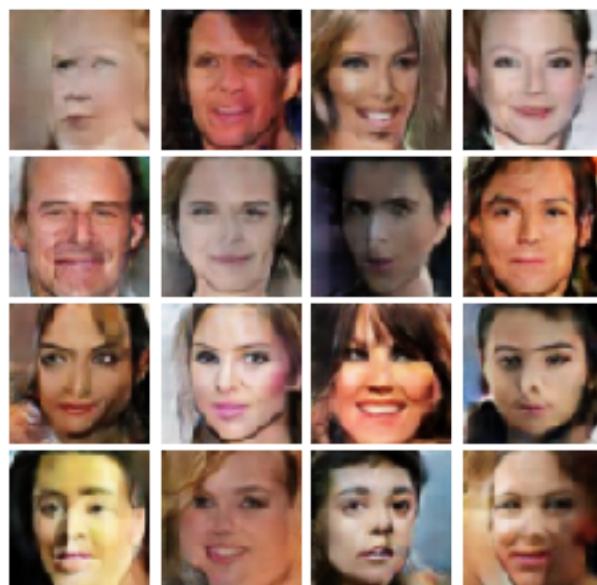
Iter: 14550, D: 0.218, G:4.549



Iter: 14700, D: 0.04625, G:5.338



Final image grid - Iter: 14790, D: 0.7842, G:6.971



[]:

5.0.2 Train LS-GAN

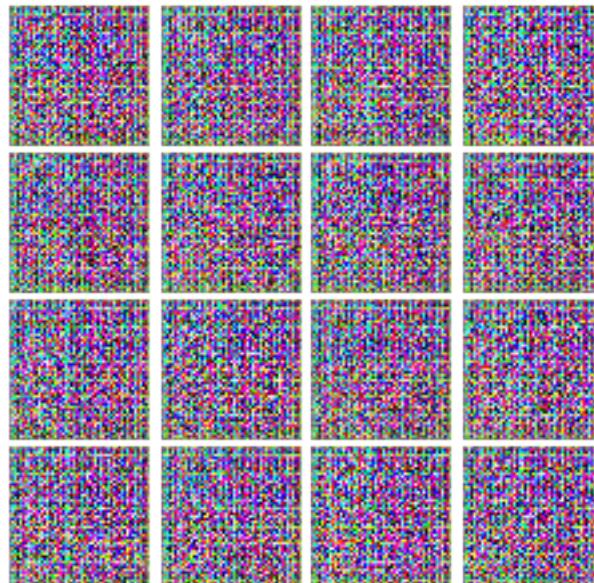
```
[ ]: D = Discriminator().to(device)
G = Generator(noise_dim=NOISE_DIM).to(device)

[ ]: D_optimizer = torch.optim.Adam(D.parameters(), lr=learning_rate, betas = (0.5, 0.999))
G_optimizer = torch.optim.Adam(G.parameters(), lr=learning_rate, betas = (0.5, 0.999))

[ ]: # ls-gan - spectral normalization
train(D, G, D_optimizer, G_optimizer, ls_discriminator_loss,
      ls_generator_loss, num_epochs=NUM_EPOCHS, show_every=200,
      train_loader=celeba_loader_train, device=device)
```

EPOCH: 1

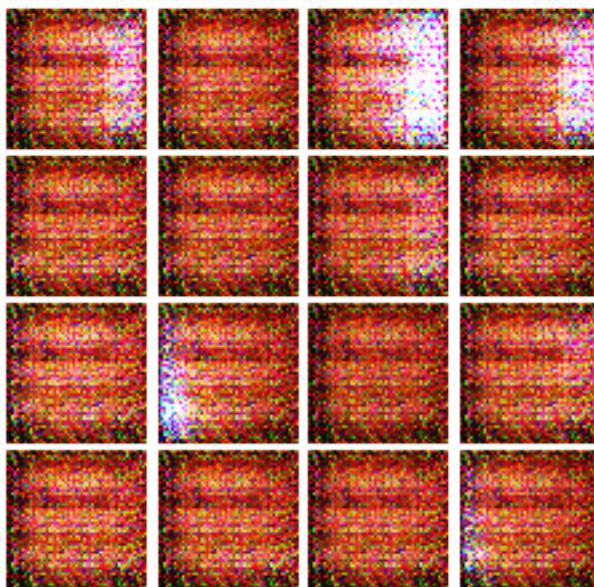
Iter: 0, D: 0.3325, G:0.1793



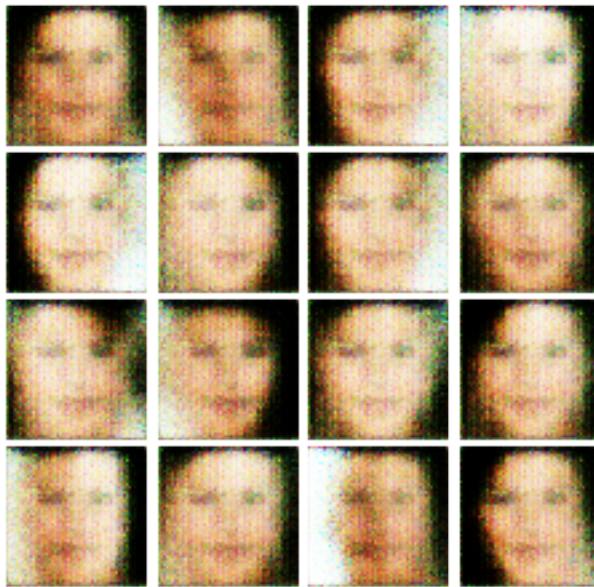
Iter: 200, D: 0.1027, G:0.5327



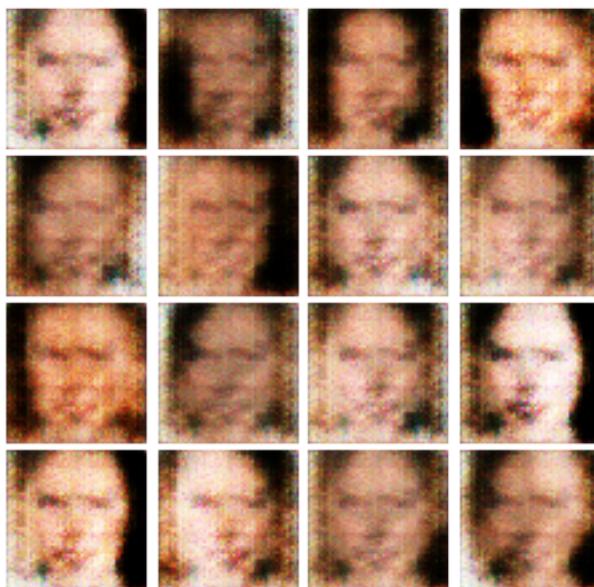
Iter: 400, D: 0.0109, G:0.5057



Iter: 600, D: 0.06717, G:0.5598

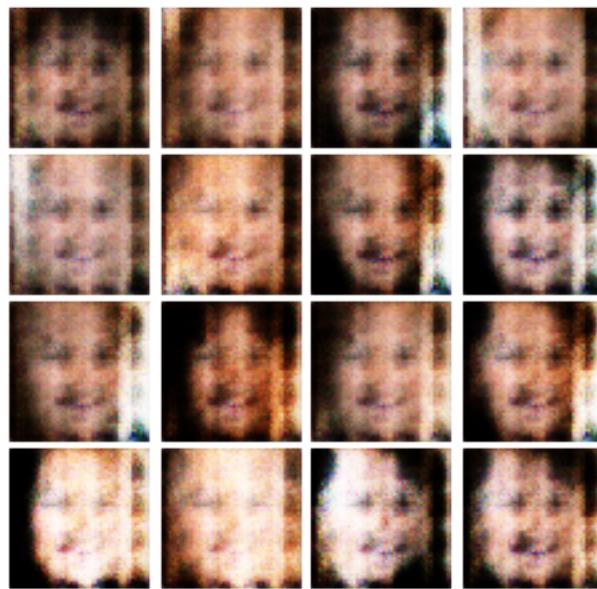


Iter: 800, D: 0.1666, G:0.5095

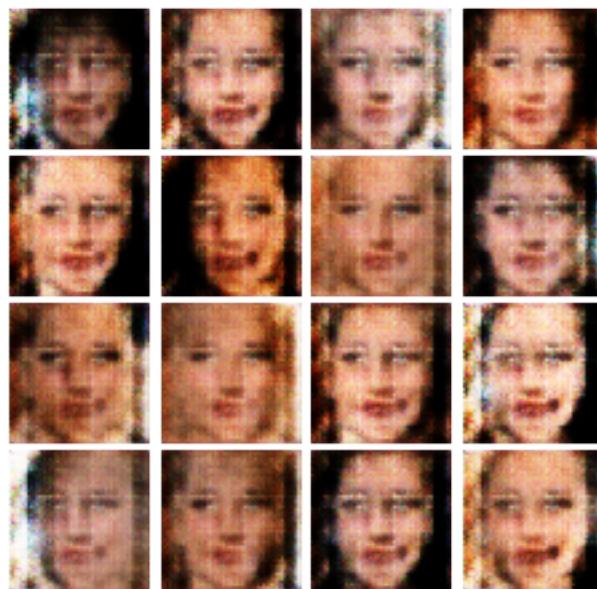


EPOCH: 2

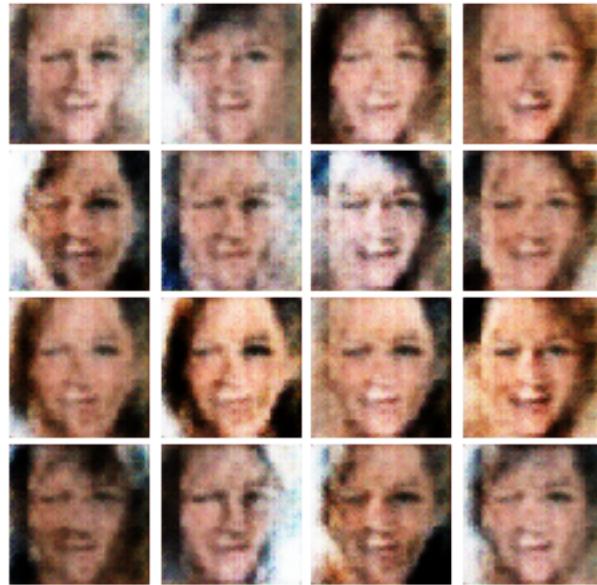
Iter: 1000, D: 0.2161, G:0.1688



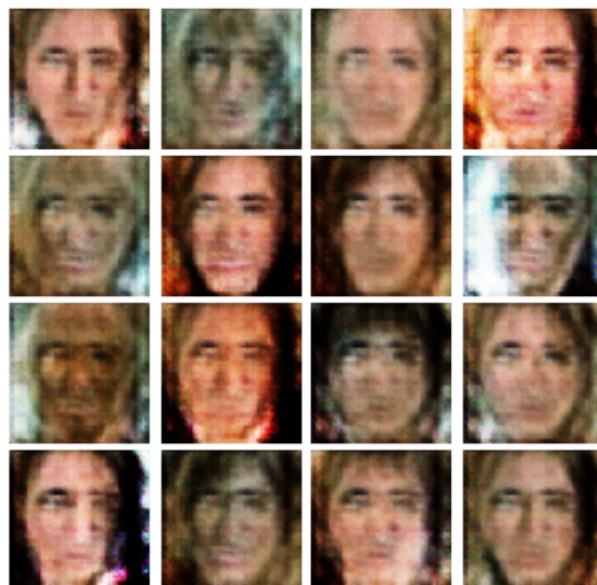
Iter: 1200, D: 0.08516, G:0.436



Iter: 1400, D: 0.08307, G:0.3749



Iter: 1600, D: 1.281, G:0.09224



Iter: 1800, D: 0.137, G:0.2429



EPOCH: 3

Iter: 2000, D: 0.2557, G: 0.2412



Iter: 2200, D: 0.3204, G: 0.5532



Iter: 2400, D: 0.1269, G:0.341



Iter: 2600, D: 0.4402, G:0.2344

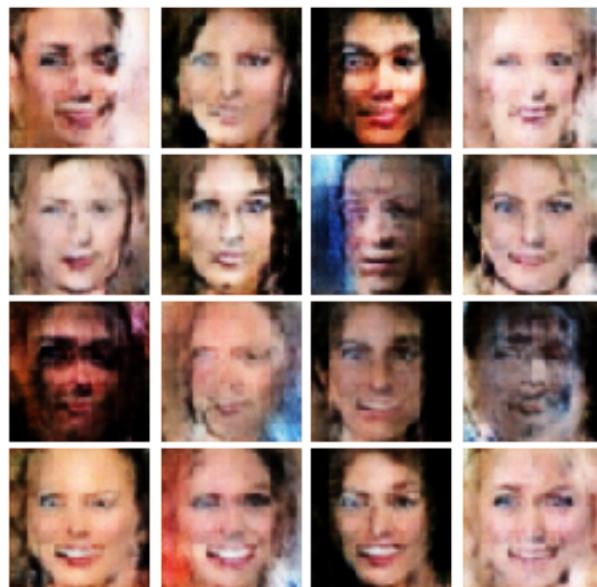


Iter: 2800, D: 0.1316, G: 0.3972

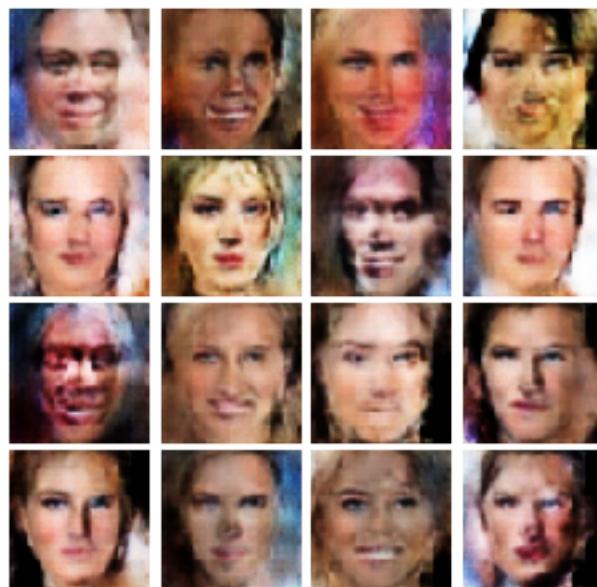


EPOCH: 4

Iter: 3000, D: 0.1446, G: 0.5463



Iter: 3200, D: 0.131, G:0.3366



Iter: 3400, D: 0.1847, G:0.1646



Iter: 3600, D: 0.1742, G: 0.4114



Iter: 3800, D: 0.1976, G: 0.4303



EPOCH: 5

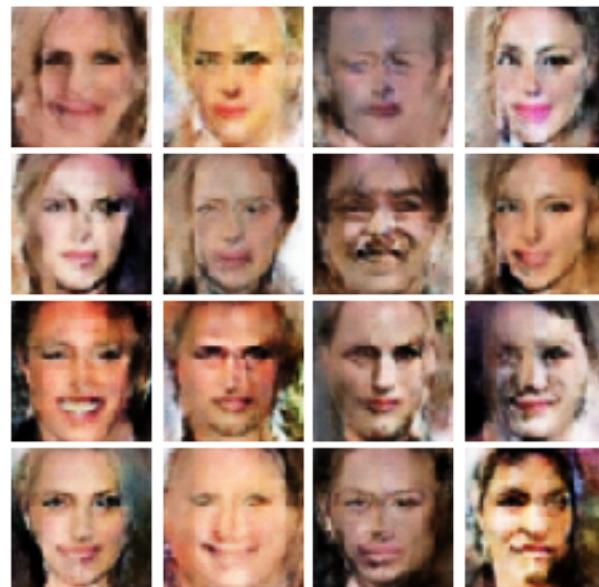
Iter: 4000, D: 0.4341, G: 0.3214



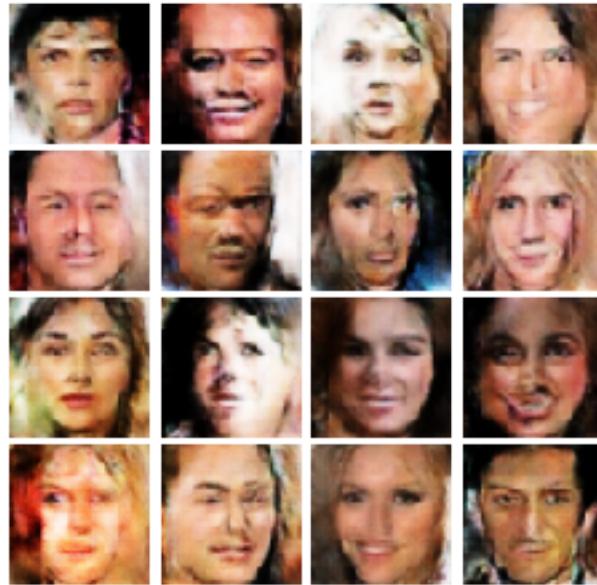
Iter: 4200, D: 0.2254, G: 0.9207



Iter: 4400, D: 0.1602, G:0.4675



Iter: 4600, D: 0.02914, G:0.485

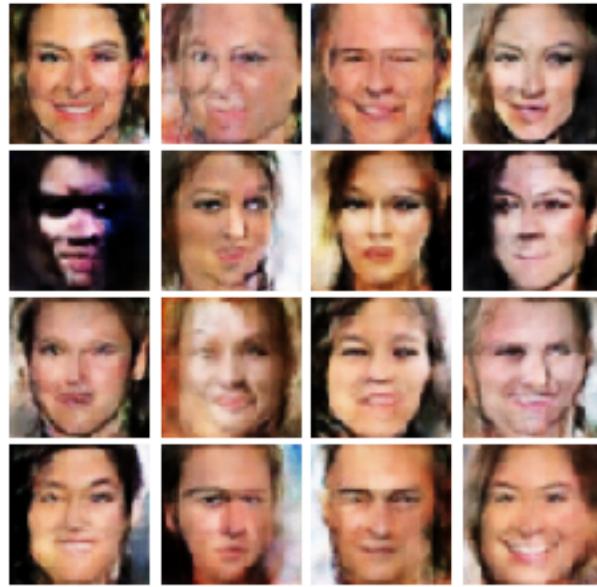


Iter: 4800, D: 0.05084, G:0.7015

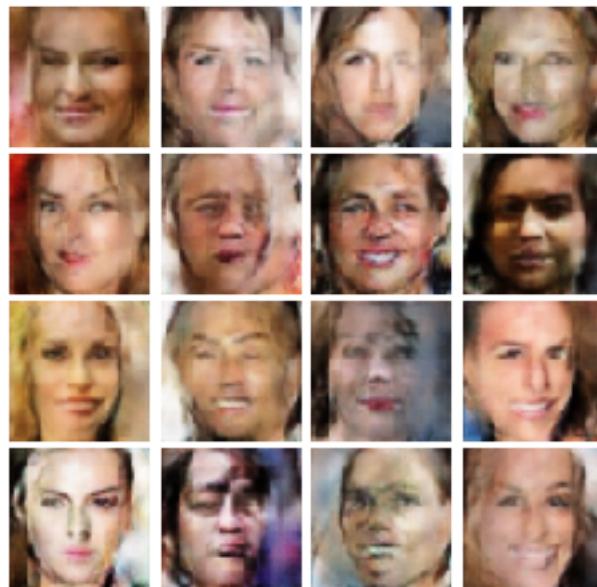


EPOCH: 6

Iter: 5000, D: 0.08078, G:0.3763



Iter: 5200, D: 0.1033, G:1.069



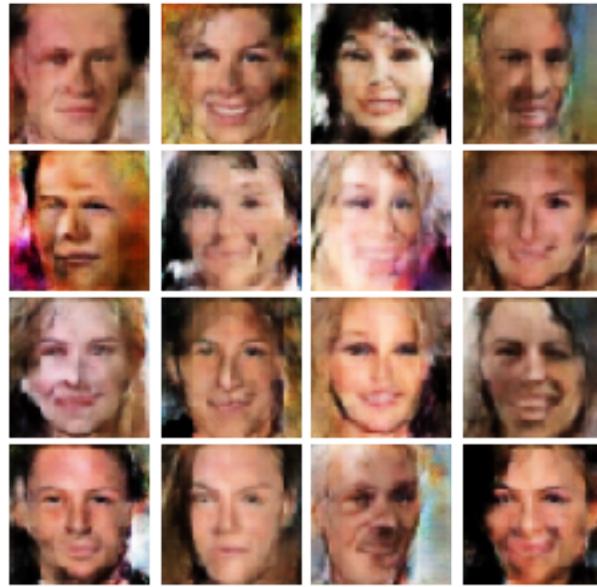
Iter: 5400, D: 0.1785, G:0.9083



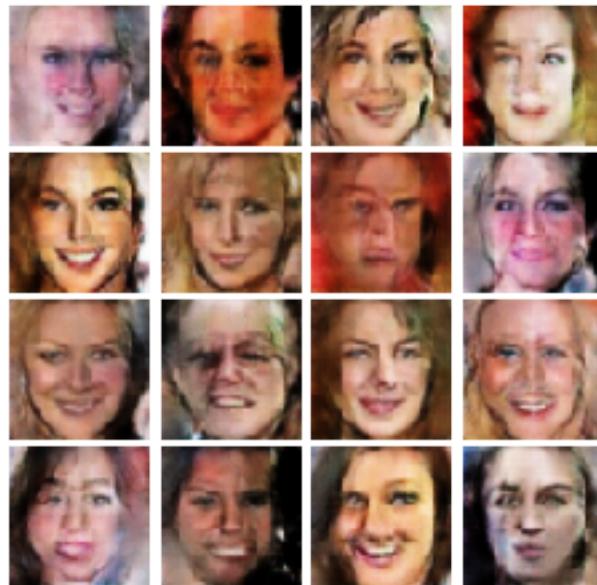
Iter: 5600, D: 0.03581, G:0.4252



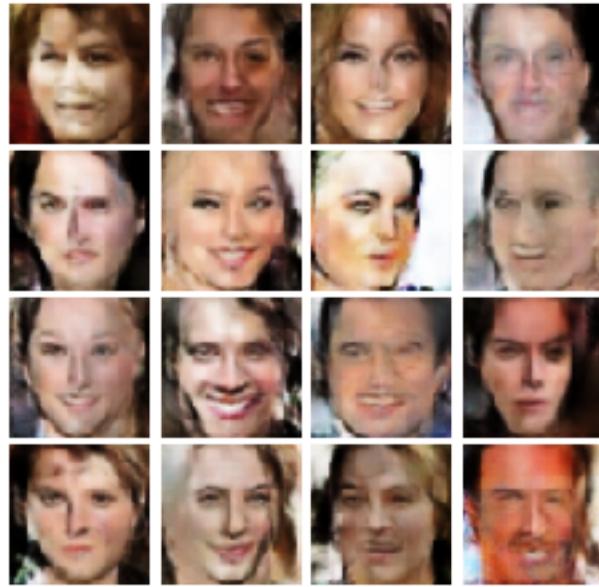
Iter: 5800, D: 0.1095, G:0.2689



EPOCH: 7
Iter: 6000, D: 0.05748, G:0.5118



Iter: 6200, D: 0.06896, G:0.4717



Iter: 6400, D: 0.05794, G:0.4825



Iter: 6600, D: 0.06395, G:0.2904

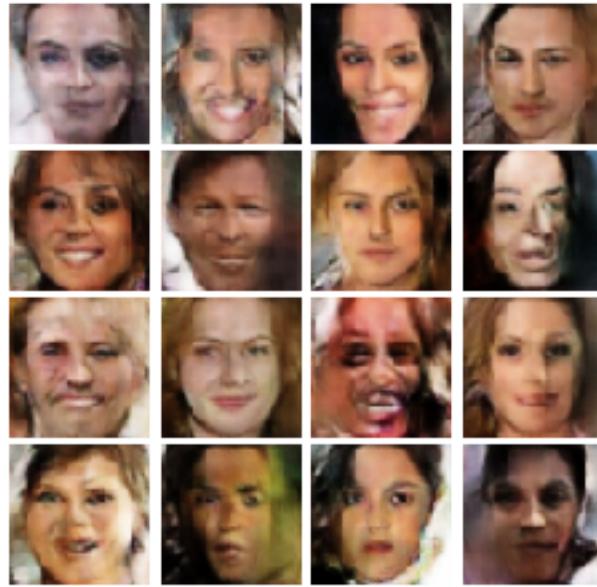


Iter: 6800, D: 0.03844, G: 0.4197

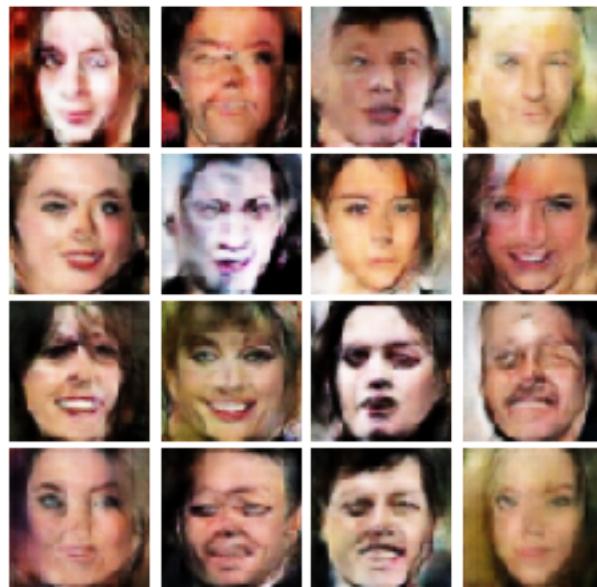


EPOCH: 8

Iter: 7000, D: 0.09331, G: 0.3998



Iter: 7200, D: 0.05645, G:0.543



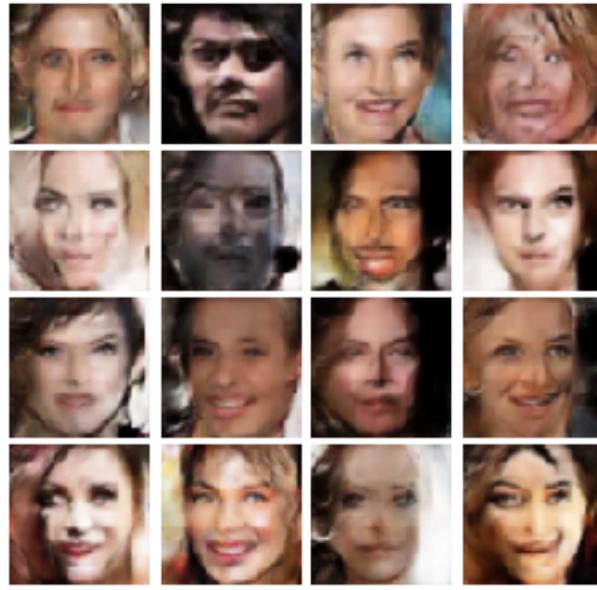
Iter: 7400, D: 0.03876, G:0.2877



Iter: 7600, D: 0.07393, G:0.3771

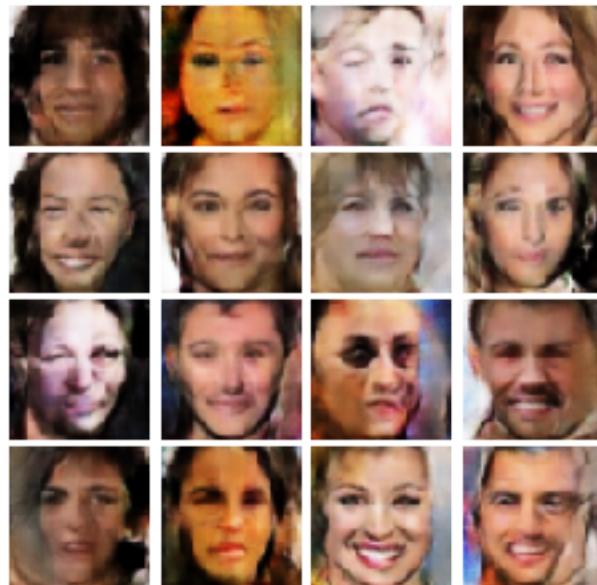


Iter: 7800, D: 0.08661, G:0.3906



EPOCH: 9

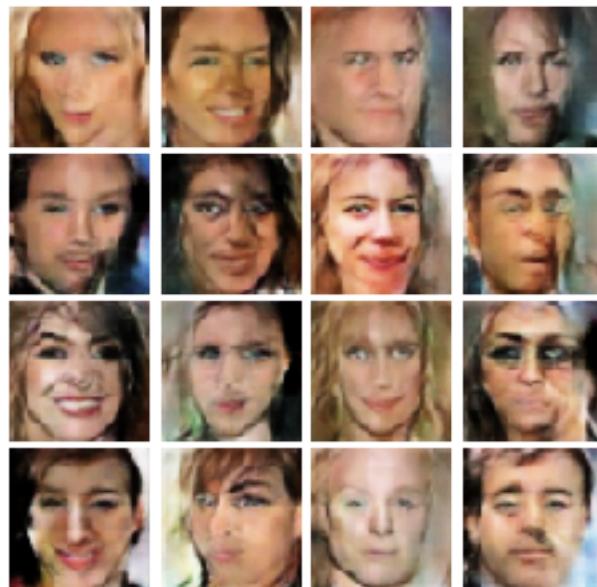
Iter: 8000, D: 0.1779, G: 0.2587



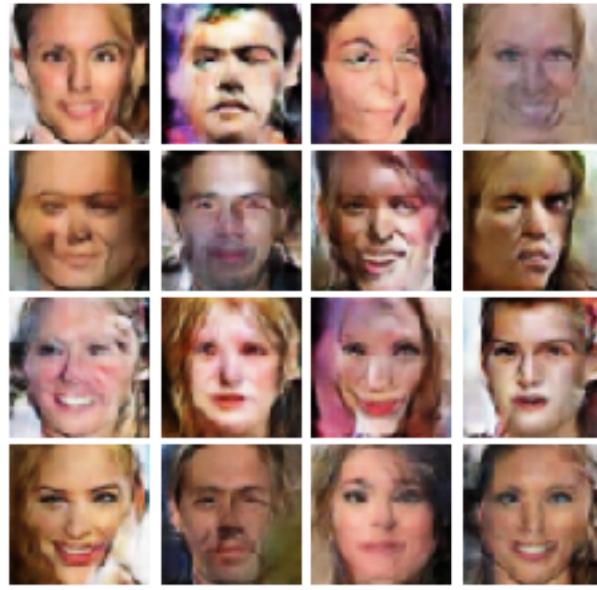
Iter: 8200, D: 0.03736, G: 0.3427



Iter: 8400, D: 0.1345, G:0.3437



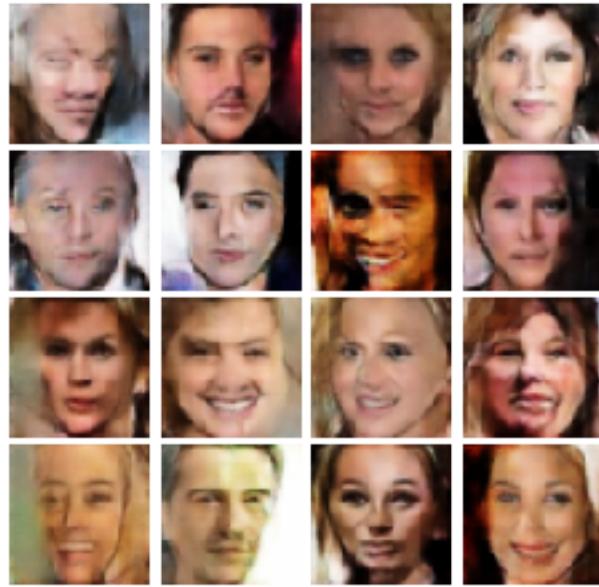
Iter: 8600, D: 0.3797, G:0.2068



Iter: 8800, D: 0.04186, G:0.7001



EPOCH: 10
Iter: 9000, D: 0.02014, G:0.4658



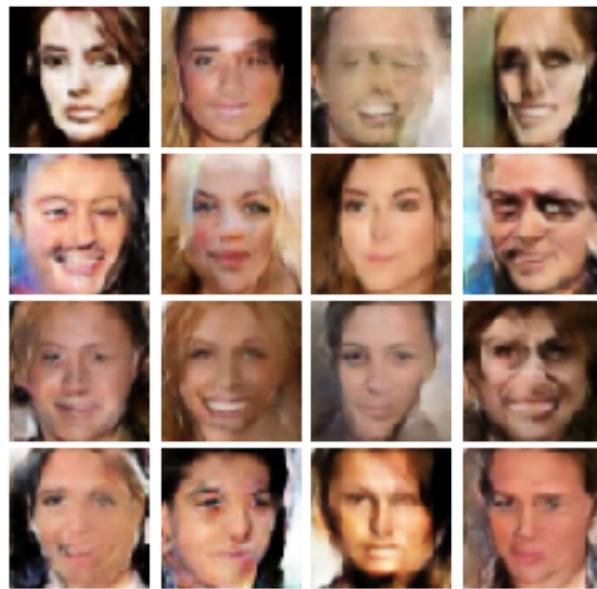
Iter: 9200, D: 0.07356, G:0.4286



Iter: 9400, D: 0.04038, G:0.5274



Iter: 9600, D: 0.02959, G:0.5609



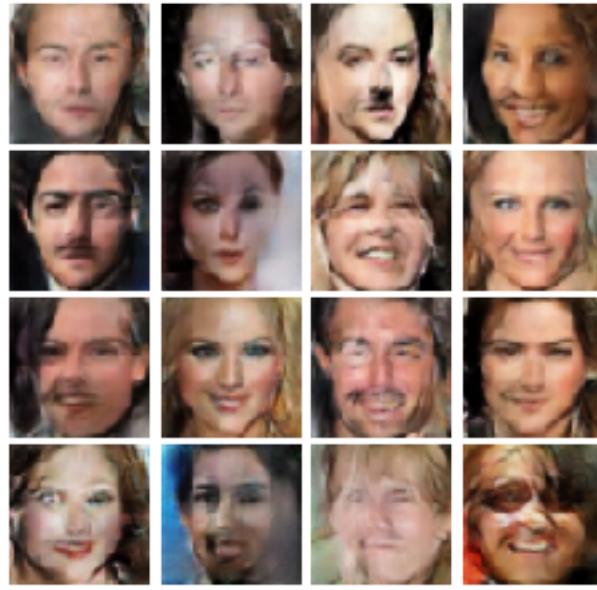
Iter: 9800, D: 0.08957, G:0.7525



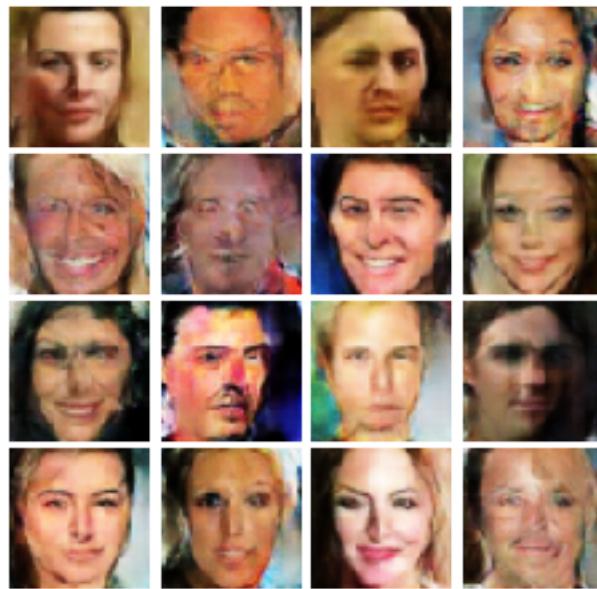
EPOCH: 11
Iter: 10000, D: 0.0333, G:0.255



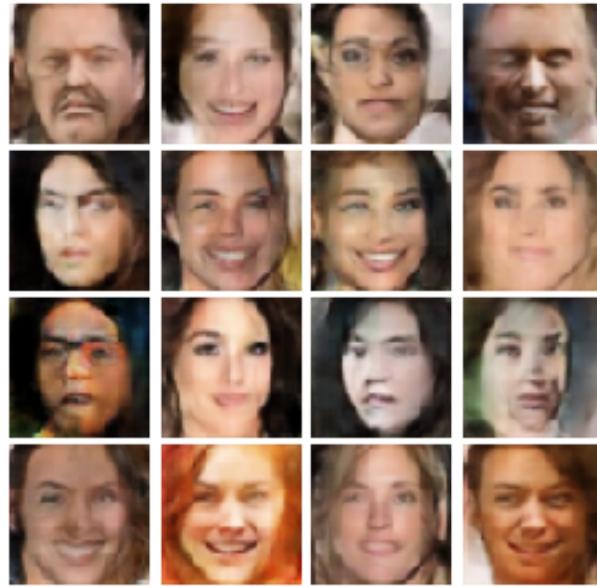
Iter: 10200, D: 0.07882, G:0.5127



Iter: 10400, D: 0.05538, G:0.6367



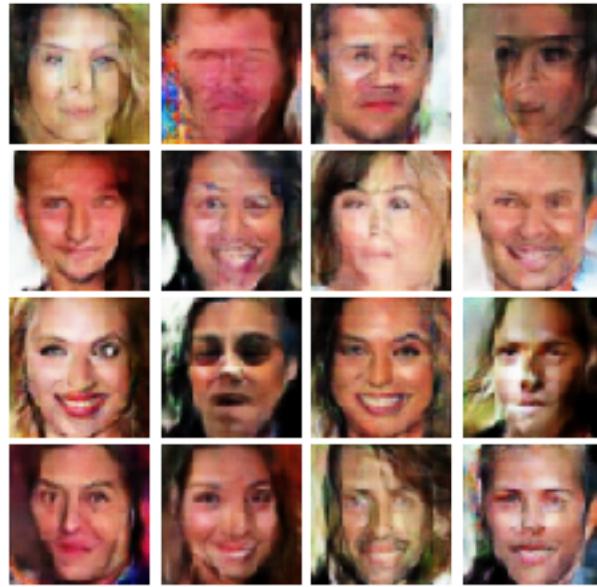
Iter: 10600, D: 0.09575, G:0.312



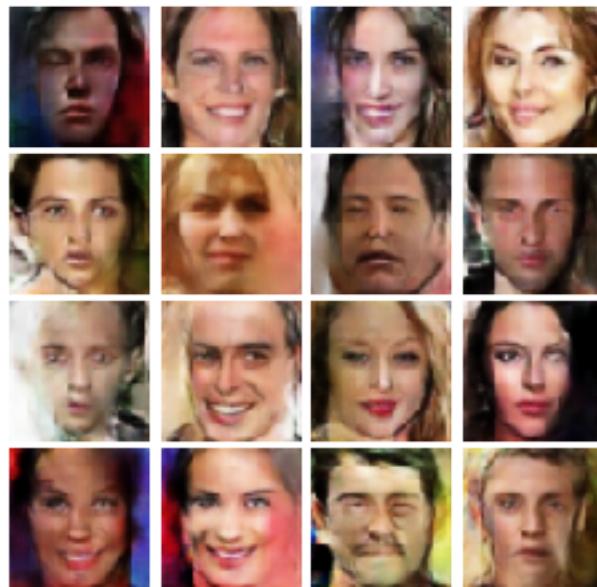
Iter: 10800, D: 0.02387, G: 0.5169



EPOCH: 12
Iter: 11000, D: 0.07836, G: 0.3941



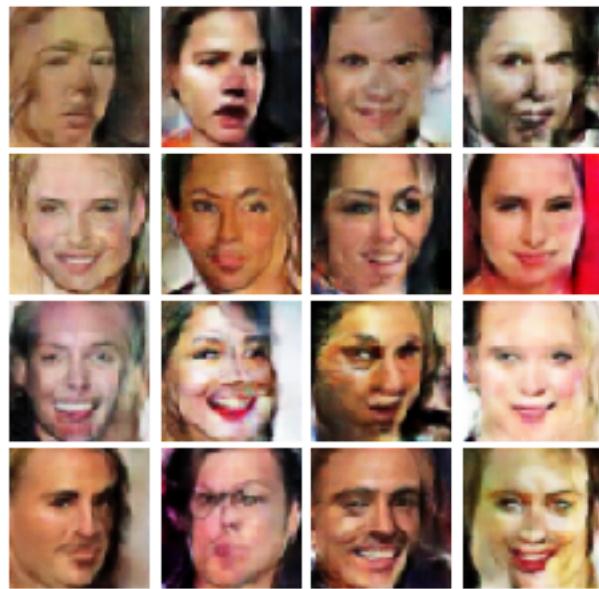
Iter: 11200, D: 0.03676, G:0.3295



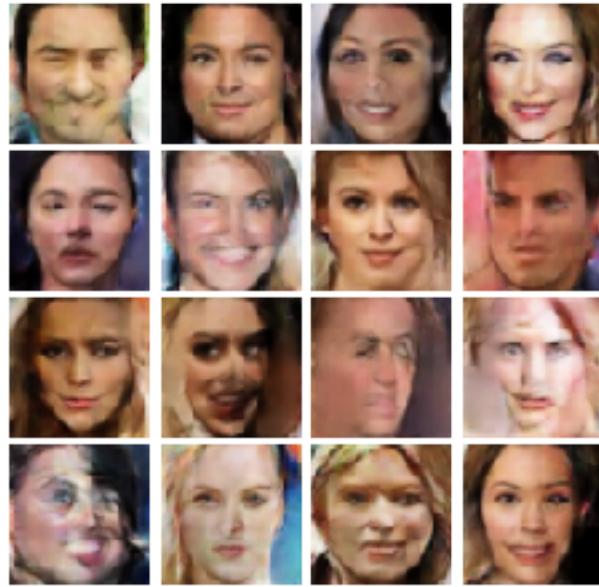
Iter: 11400, D: 0.0429, G:0.3114



Iter: 11600, D: 0.0984, G:0.7645

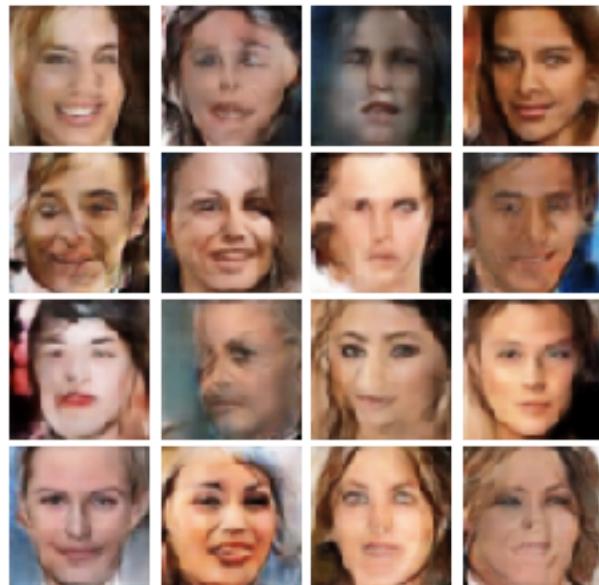


Iter: 11800, D: 0.03191, G:0.4425

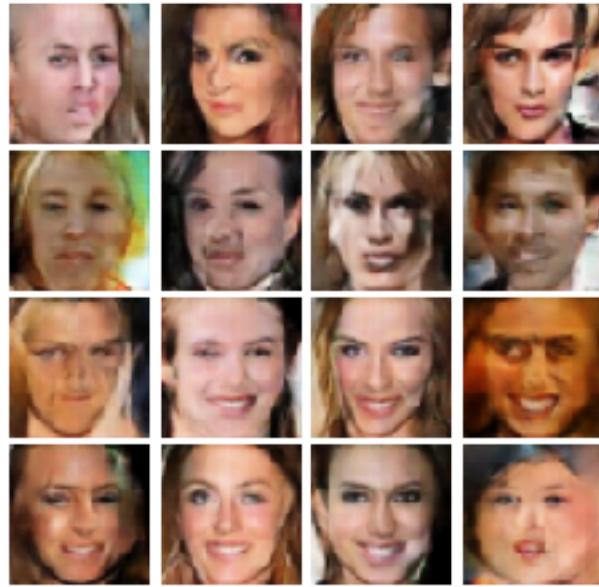


EPOCH: 13

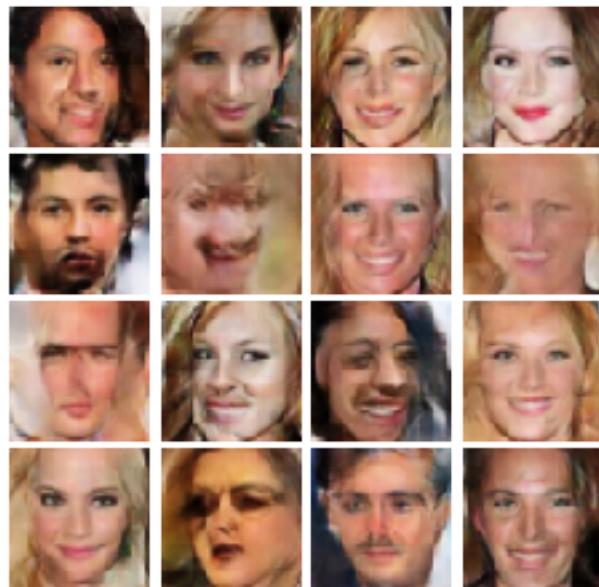
Iter: 12000, D: 0.0214, G:0.5287



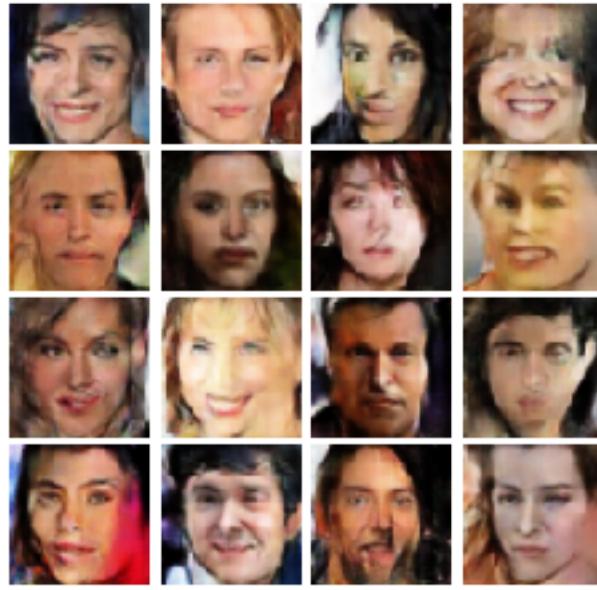
Iter: 12200, D: 0.0621, G:0.4768



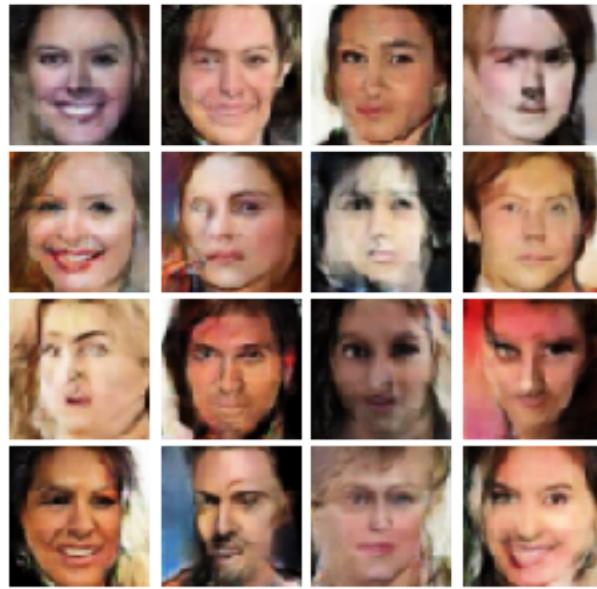
Iter: 12400, D: 0.07967, G:0.4969



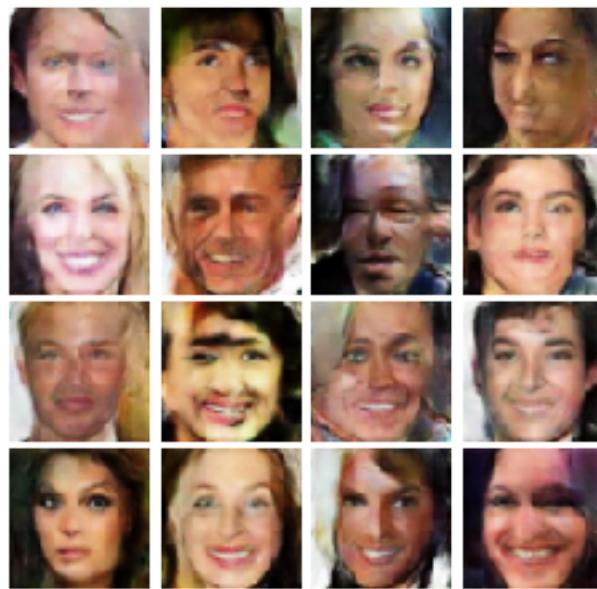
Iter: 12600, D: 0.05732, G:0.7469



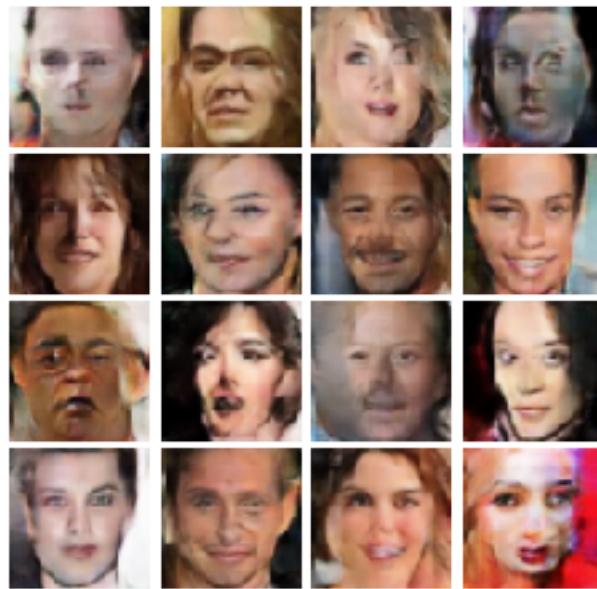
Iter: 12800, D: 0.03189, G:0.2925



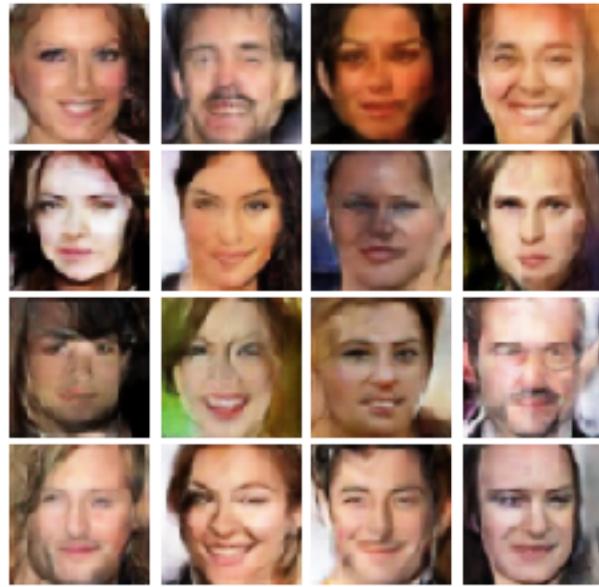
EPOCH: 14
Iter: 13000, D: 0.02664, G:0.4641



Iter: 13200, D: 0.03719, G:0.3895



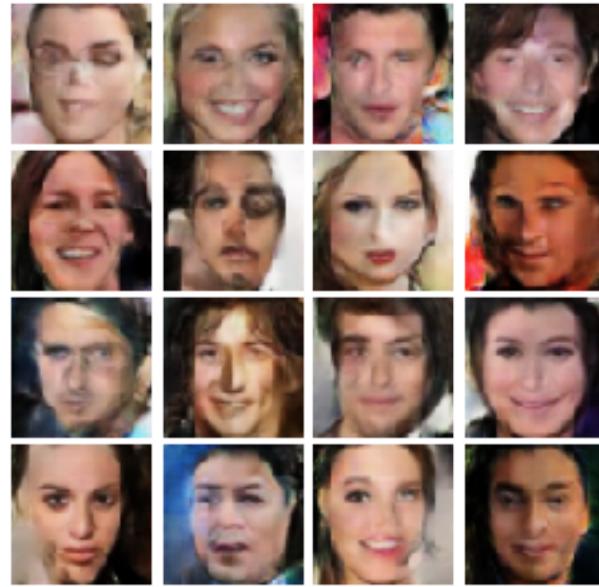
Iter: 13400, D: 0.07297, G:0.5169



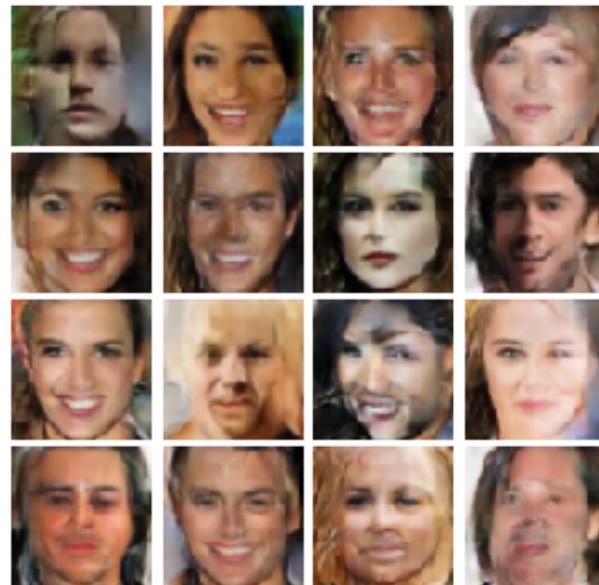
Iter: 13600, D: 0.09203, G:0.6784



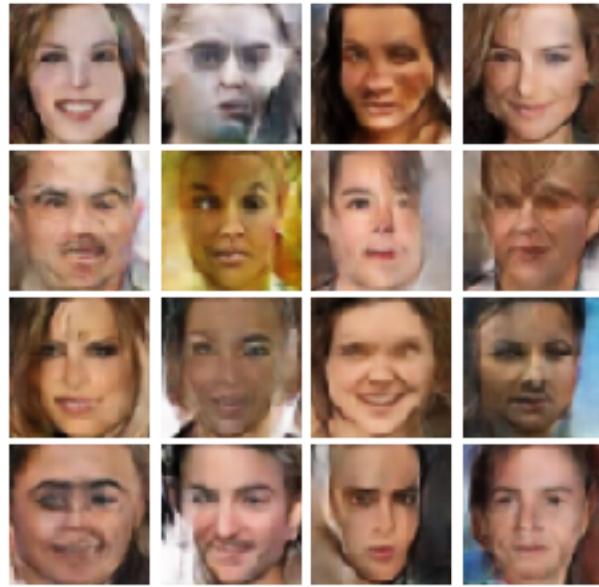
Iter: 13800, D: 0.0642, G:0.3263



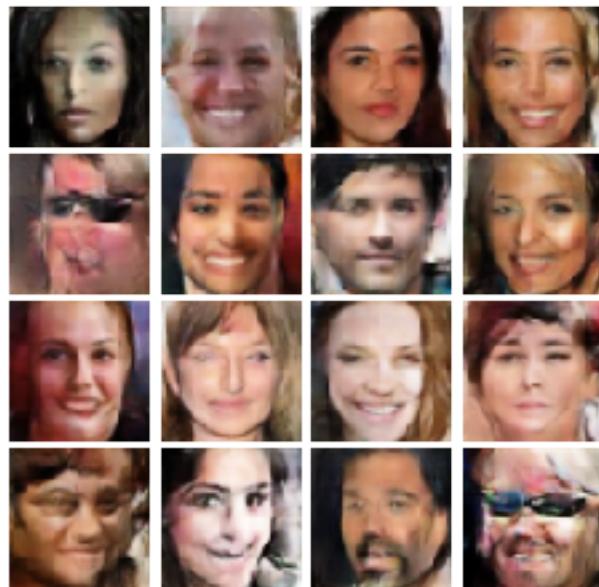
EPOCH: 15
Iter: 14000, D: 0.05362, G:0.4279



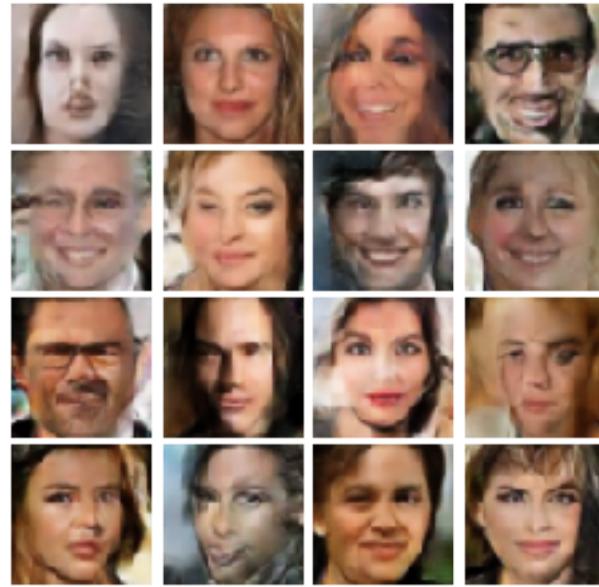
Iter: 14200, D: 0.04428, G:0.4293



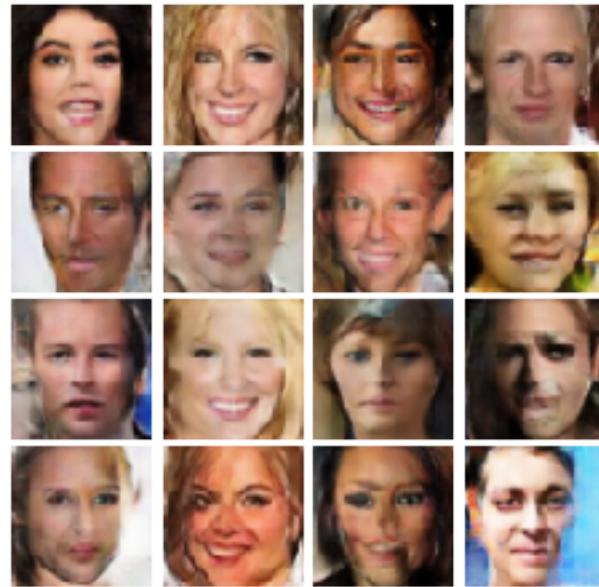
Iter: 14400, D: 0.04158, G:0.5897



Iter: 14600, D: 0.02051, G:0.4097



Final image grid - Iter: 14790, D: 0.02369, G: 0.381



[]: